

GAME ENGINE

Source

- The Cherno
- GAMES 104: Modern Game Engine – Theory and Practice

Framework

- Basic Elements
 - Structure and layer...
 - Rendering
 - Animation
 - Physics
 - Gameplay
 - Event system, scripts system, graph driven
 - Misc. Systems
 - Effects, navigation, camera...
 - Tool set
 - C++ reflection, data scheme (reflection: complex)
 - Online gaming
 - Synchronization, consistency
- Advanced tech:
Motion matching
Procedural content generation (PCG)
- Data-oriented programming (DOP)
Job system
- (UE5 amazing systems)
Lumen
Nanite

Layered Architecture of Game Engine

- Tool Layer (chain of editors)
 - Function Layer (make it visible, movable and playable)
 - Resource Layer (data and file)
 - Core Layer (swiss knife of game engine)
 - Platform Layer (launch on different platforms)
- +Middleware and 3rd party libraries

Why:

- Decoupling and Reducing Complexity
- Response for Evolving Demands

Resource

how to access my data

Offline Resource Importing

- Unify file access by defining a meta asset file format (ie.ast)
- Assets are faster to access by importing preprocess
- Build a composite asset file to refer to all resources
- GUID is an extra protection of reference

Manage asset life cycle

Memory management for Resources - life cycle

- Different resources have different life cycles
- Limited memory requires release of loaded resources when possible
- Garbage collection and deferred loading is critical features

Resources (Game Assets)

3D Model
Resource

Texture
Resource

Material
Resource

Font
Resource

Skeleton
Resource

Collision
Resource

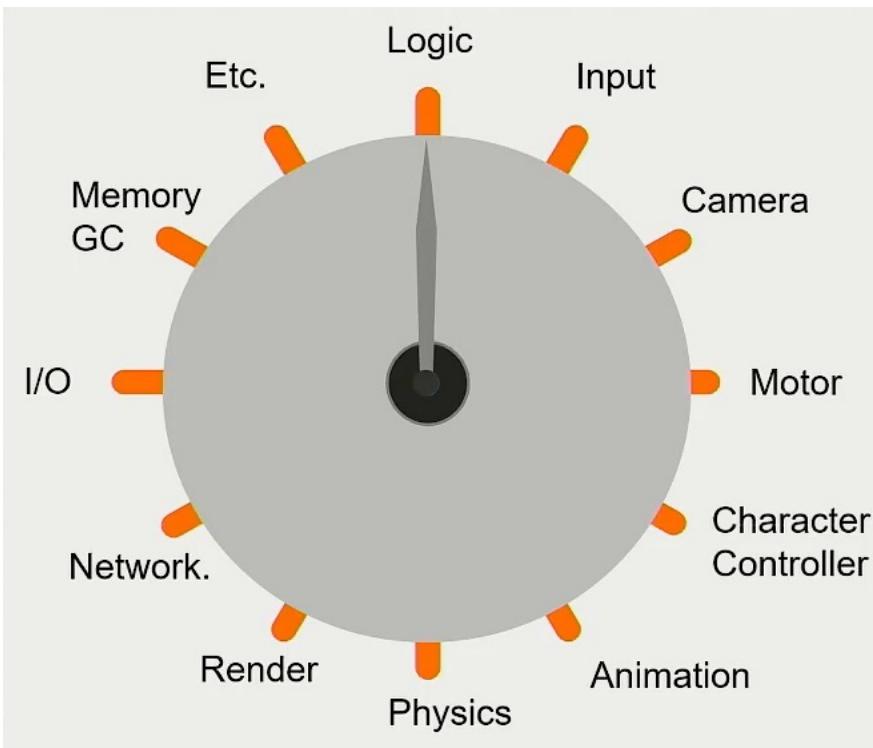
Physics
Parameters

Game
World/Map

etc.

Resource Manager

Function



Tick
--logic
--render

```
void tickMain(float delta_time)
{
    while (!exit_flag)
    {
        tickLogic(delta_time);
        tickRender(delta_time);
    }
}
```

```
void tickLogic(float delta_time)
{
    tickCamera(delta_time);
    tickMotor(delta_time);
    tickController(delta_time);
    tickAnimation(delta_time);
    tickPhysics(delta_time);
    /*...*/
}
```

```
void tickRender(float delta_time)
{
    tickRenderCamera();
    culling();
    rendering();
    postprocess();
    present();
}
```

Heavy-duty hotchpotch
Multi-threading

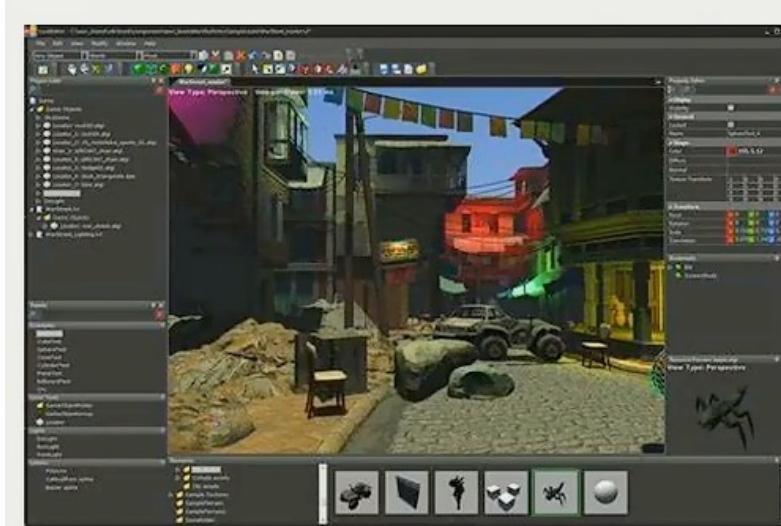
Core

- Math Library
- Math Efficiency(quick and dirty hacks, SIMD)
- Data structure and containers (customized STL)
- Memory management (memory pool).

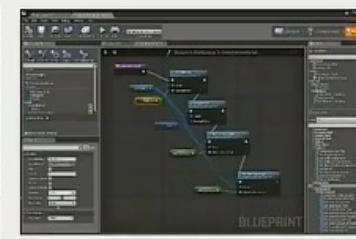
Platform

- File system
- Graphics API (DirectX, Vulkan...)
- Hardware architecture

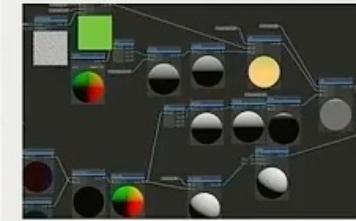
Tool



Level Editor



Logical Editor



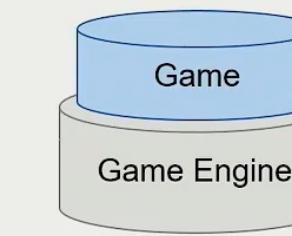
Shader Editor



Animation Editor



UI Editor



- Digital content creation (DCC)

Allow everyone to create game

Unleash the Creativity

- Build upon game engine
- Create, edit and exchange game play assets

Flexible of coding languages



Editors



DCC



Houdini



MAYA



blender

How to Build a Game World

- Game Object (GO)

- Name, property, behavior.
- Inheritance -> component base

Component

- Code example

Base class of component

```
class ComponentBase
{
    virtual void tick() = 0;
    ...
};
```

```
class GameObjectBase
{
    vector<ComponentBase*> components;
    virtual void tick();
    ...
};

class Drone: public GameObjectBase
{
    ...
};
```



```
class TransformComponent: public ComponentBase
```

```
{    Vector3 position;
    ...
    void tick();
};
```

```
class ModelComponent: public ComponentBase
```

```
{    Mesh mesh;
    ...
    void tick();
};
```

```
class MotorComponent: public ComponentBase
```

```
{    float battery;
    void tick();
    void move();
    ...
};
```

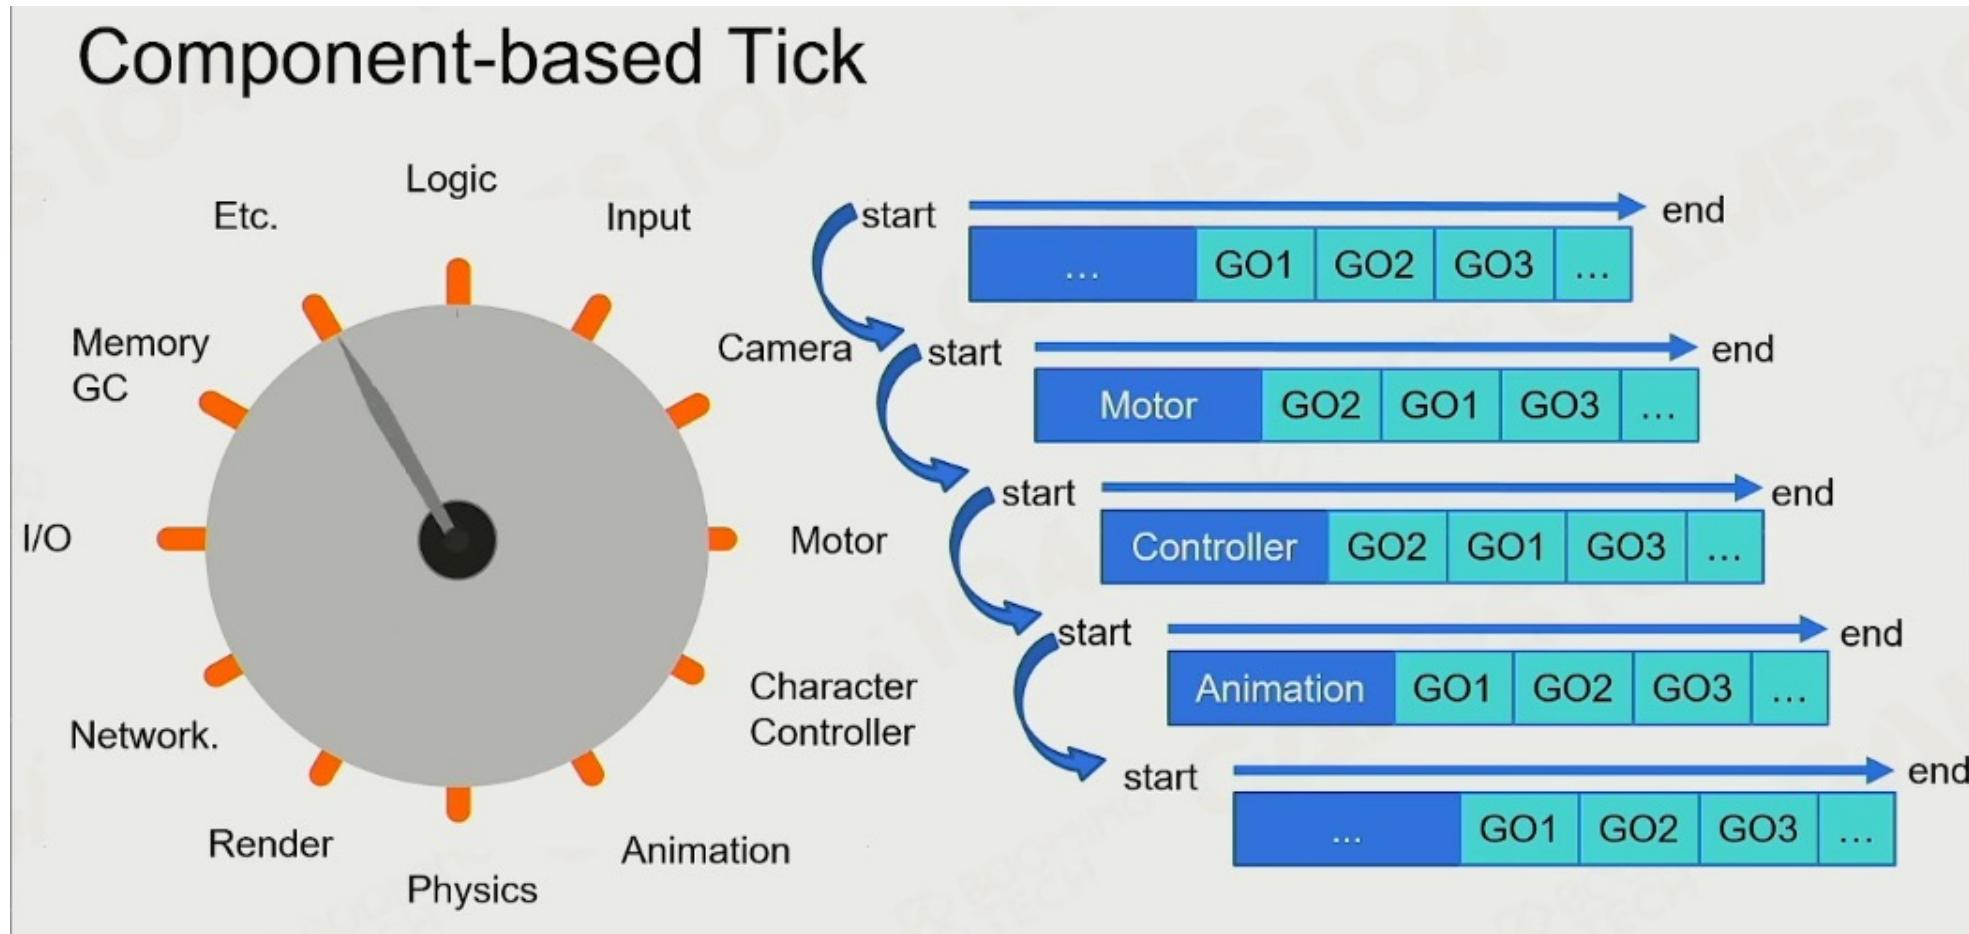
```
class AIComponent: public ComponentBase
```

```
{    void tick();
    void scout();
    ...
};
```

Animation
Physics

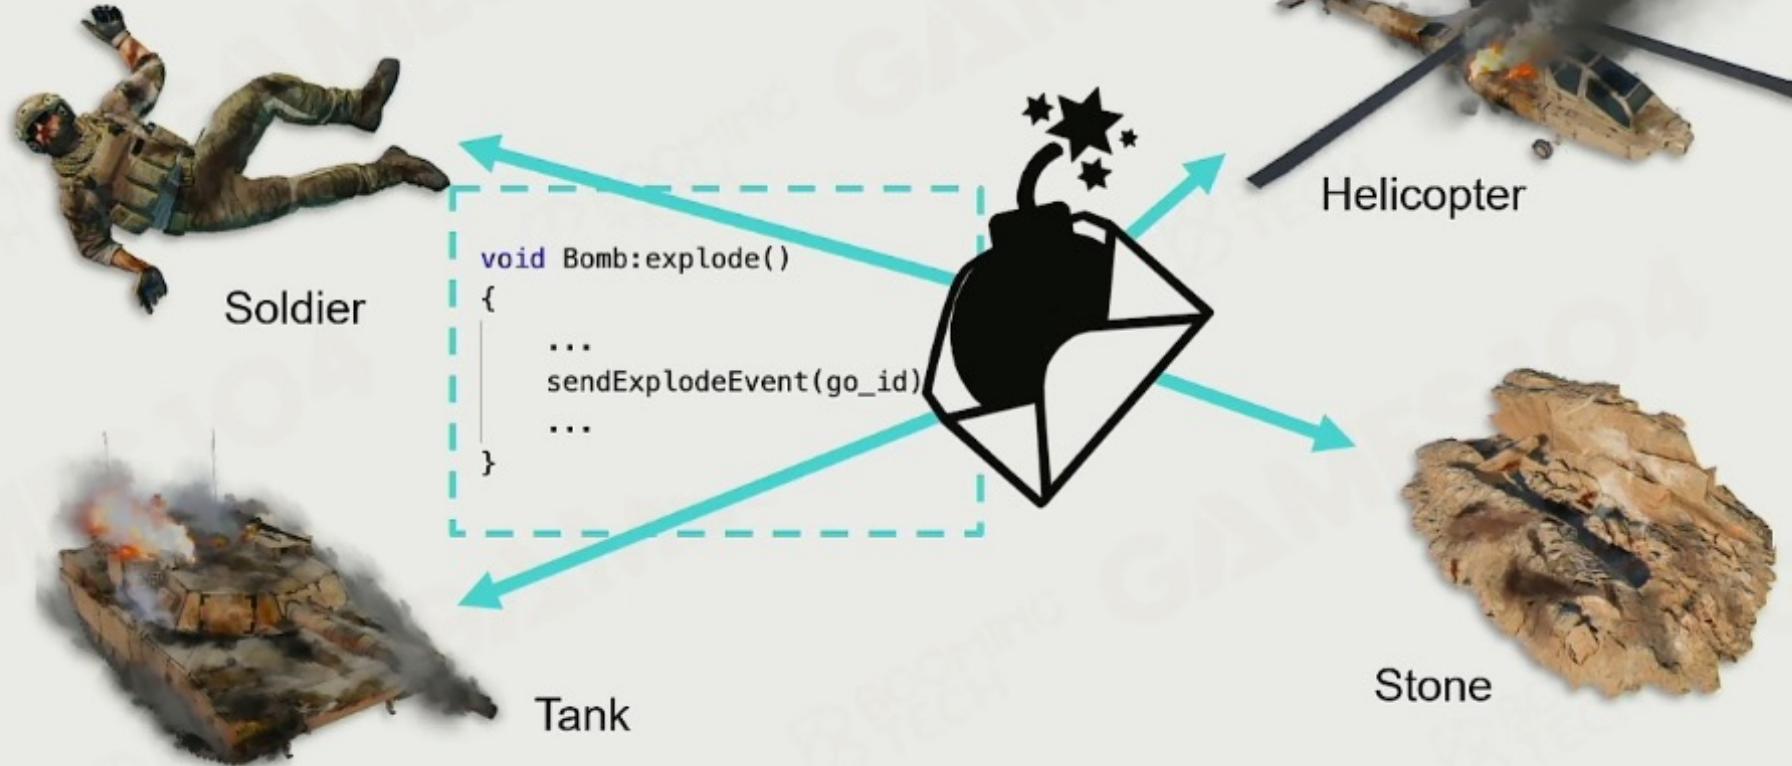


Pipeline – batch process



Events

- Message sending and handling
- Decoupling event sending and handling



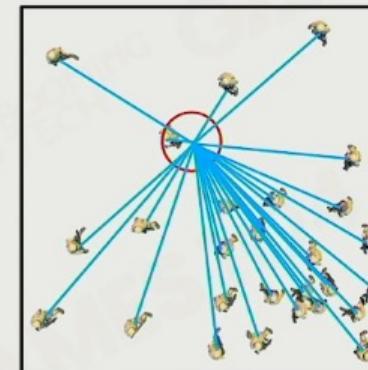
Scene Management

- Game objects are managed in a scene
- Game object query
 - By unique game object ID
 - By object position

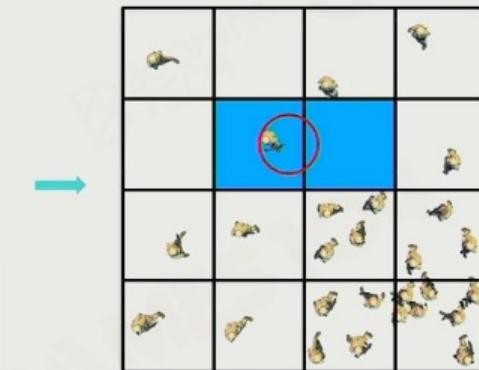


Scene Management

- Simple space segmentation



No division



Divided by grid



A needle in a haystack!

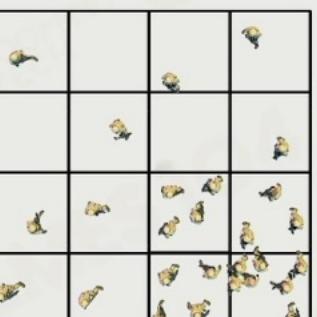


Scene Management

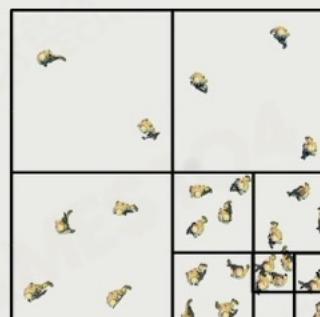
- Segmented space by object clusters
- Hierarchical segmentation



Hangzhou,
Zhejiang, China



Divided by grid

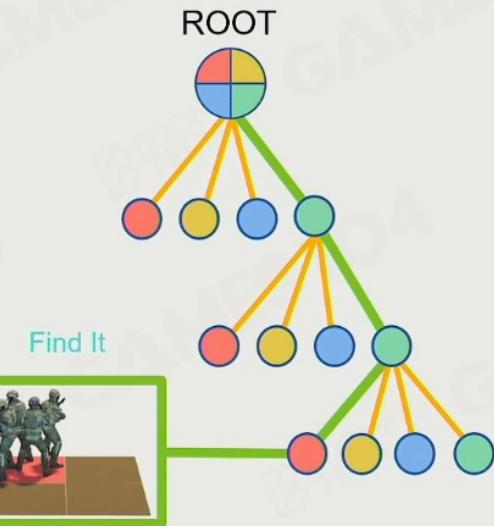
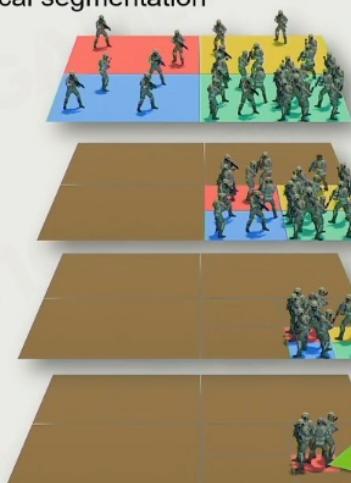


Quadtree



Scene Management

- Segmented space by object clusters
- Hierarchical segmentation

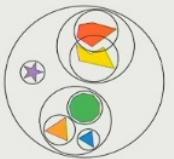


Find It

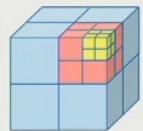
Bounding box

Scene Management

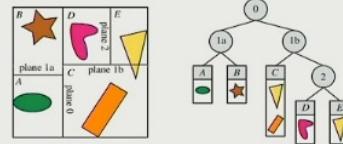
- Spatial Data Structures



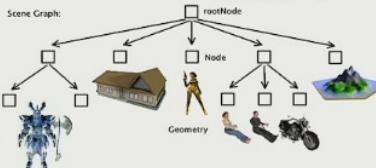
Bounding Volume Hierarchies (BVH)



Octree

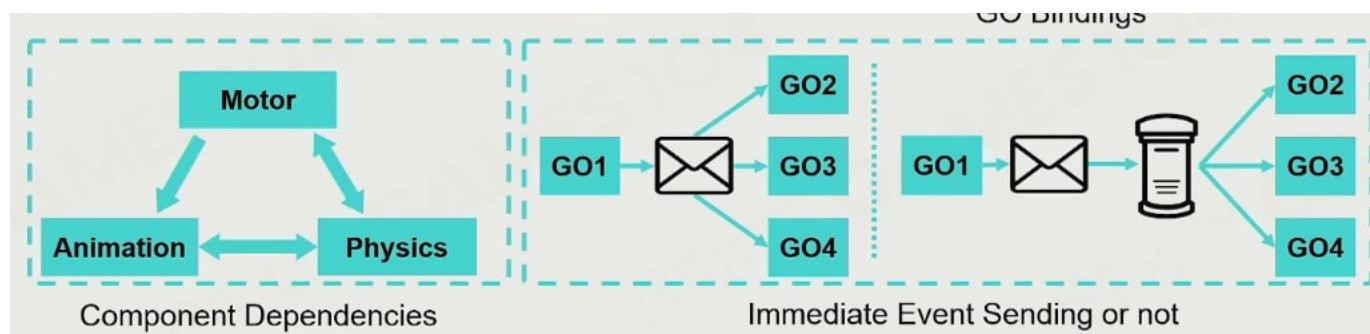


Binary Space Partitioning(BSP)



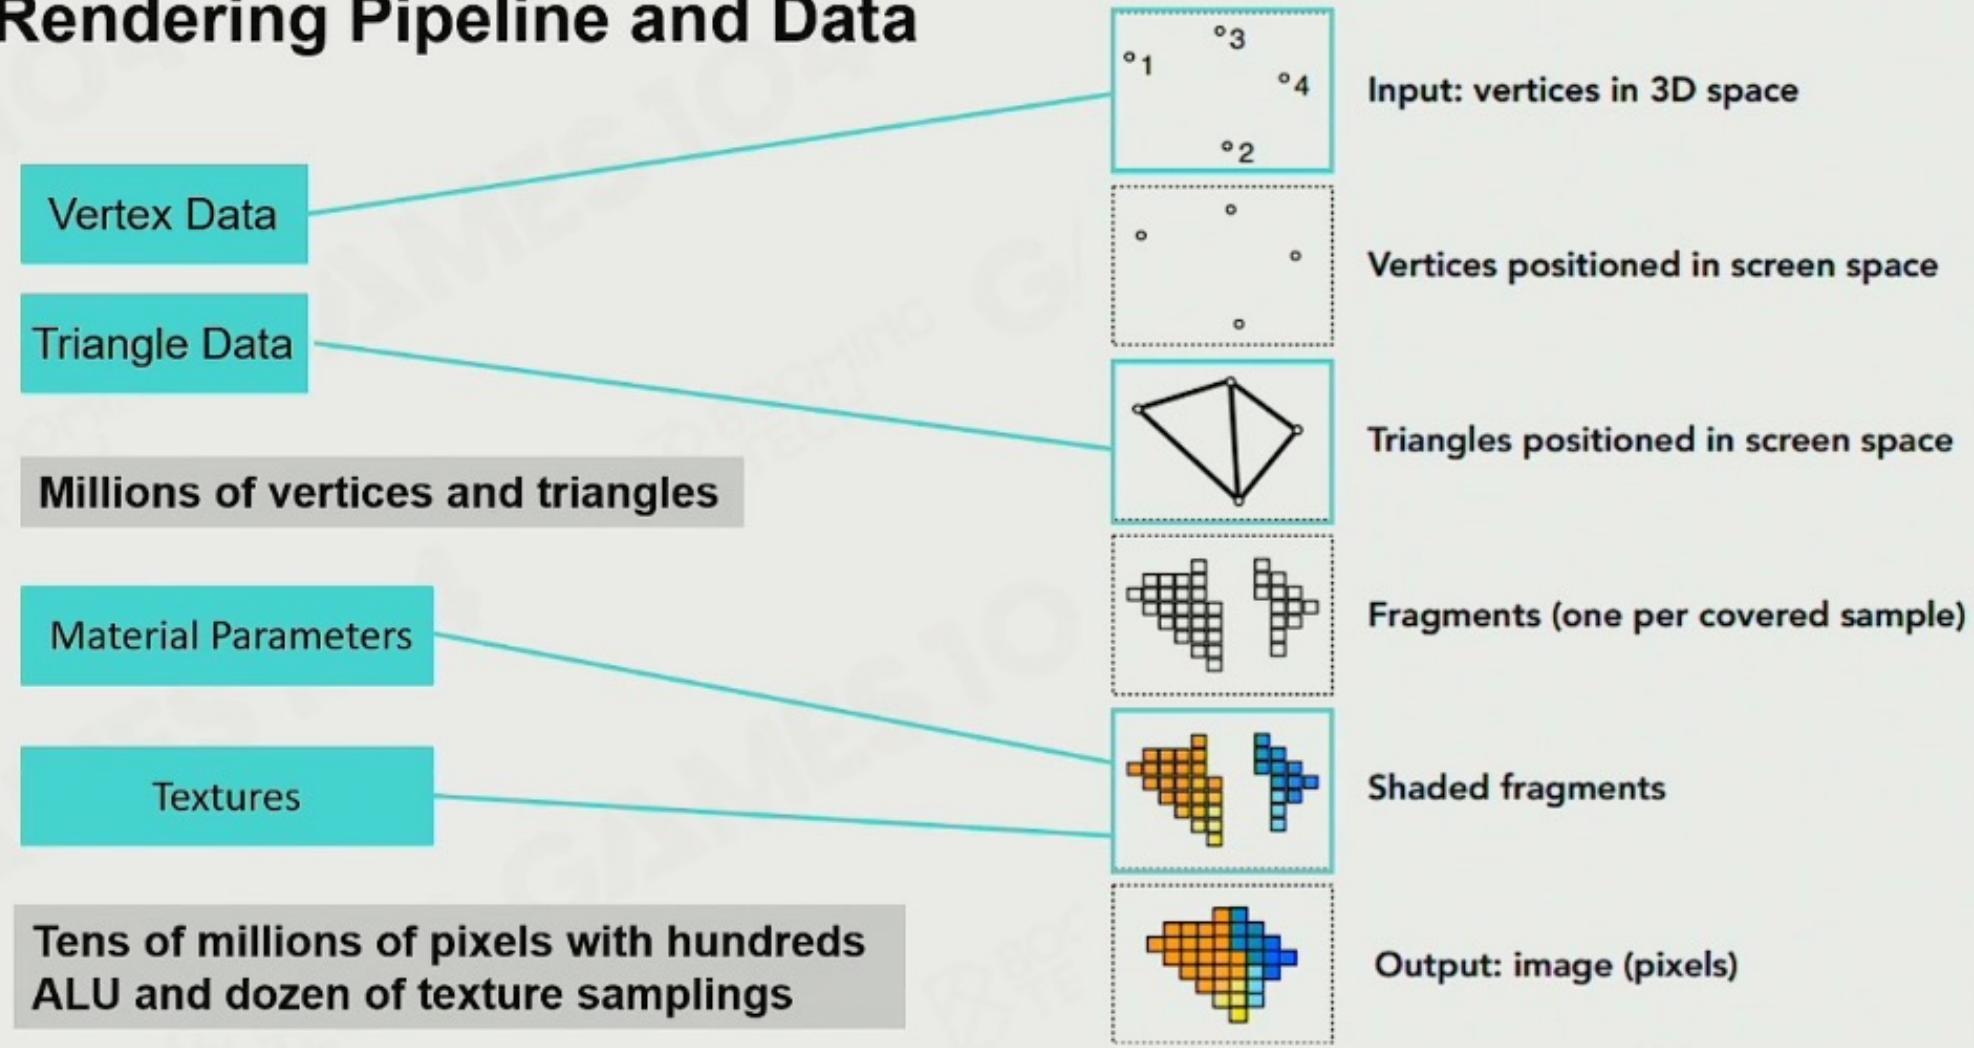
Scene Graph

parallel

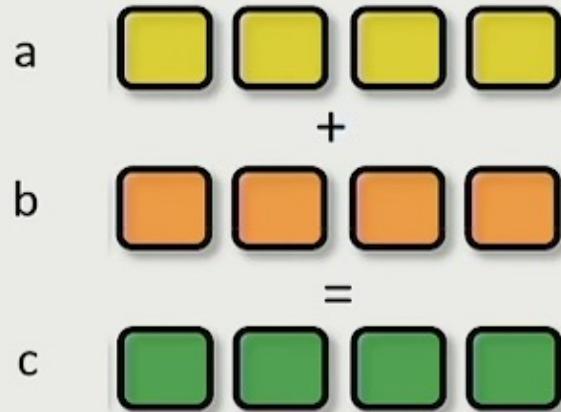


Rendering on Game Engine

Rendering Pipeline and Data



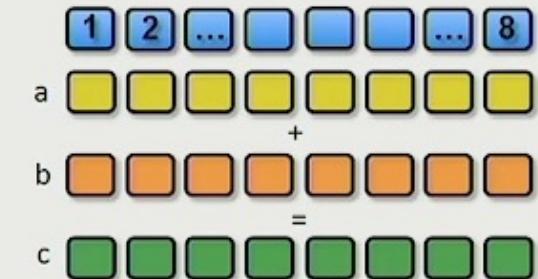
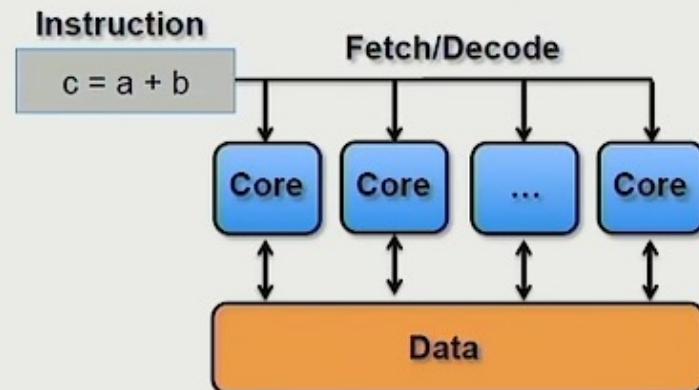
SIMD and SIMT



`SIMD_ADD c, a, b`

SIMD (Single Instruction Multiple Data)

- Describes computers with multiple processing elements that perform the same operation on multiple data points simultaneously

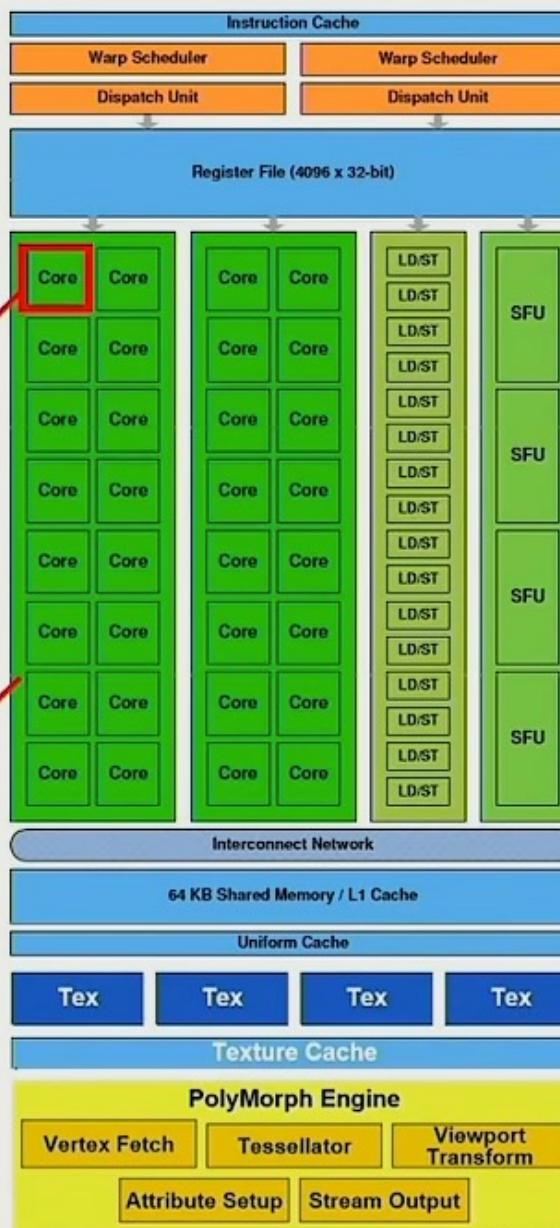


`SIMT_ADD c, a, b`

SIMT (Single Instruction Multiple Threads)

- An execution model used in parallel computing where single instruction, multiple data (SIMD) is combined with multithreading

GPU Architecture



GPC (Graphics Processing Cluster)

A dedicated hardware block for computing, rasterization, shading, and texturing

SM (Streaming Multiprocessor)

Part of the GPU that runs CUDA kernels

Texture Units

A texture processing unit, that can fetch and filter a texture

CUDA Core

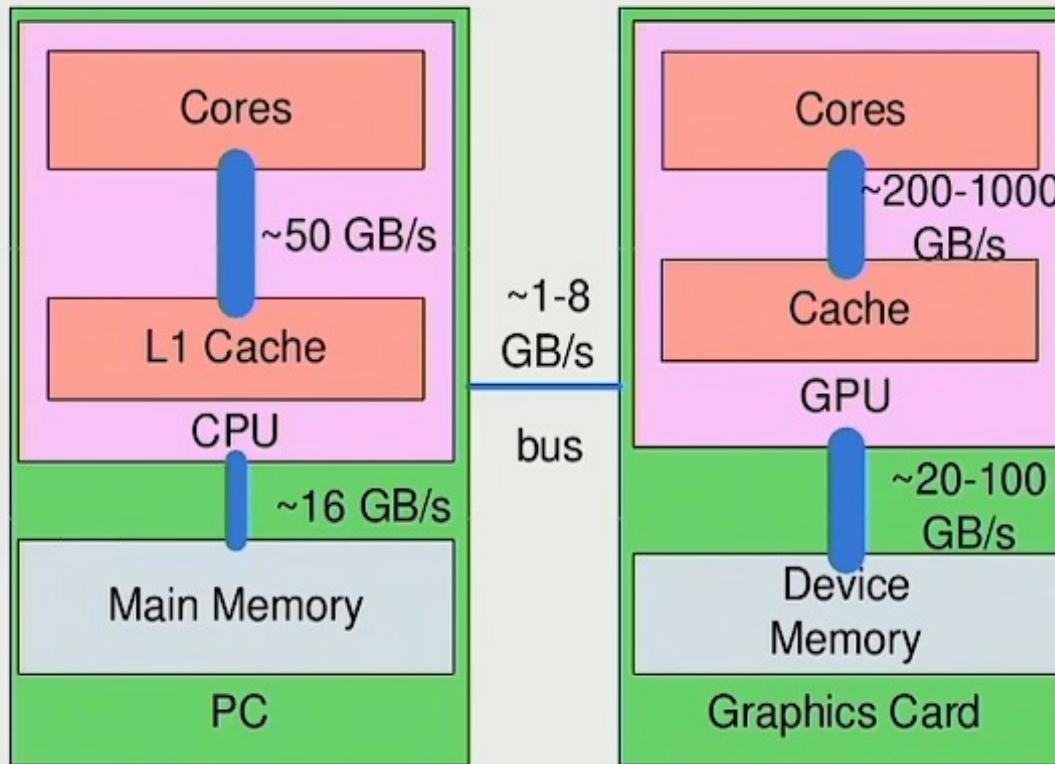
Parallel processor that allow data to be worked on simultaneously by different processors

Warp

A collection of threads

e.g. Fermi Architecture

Data Flow from CPU to GPU



- **CPU and Main Memory**
 - Data Load / Unload
 - Data Preparation
- **CPU to GPU**
 - High Latency
 - Limited Bandwidth
- **GPU and Video Memory**
 - High Performance Parallel Rendering

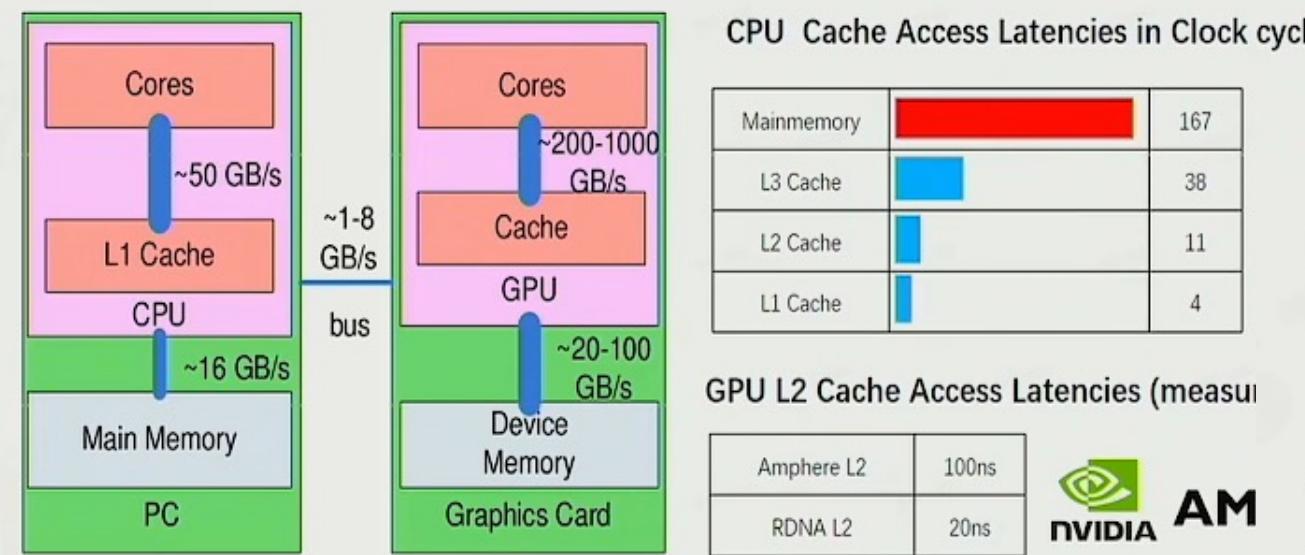
Tips

Always minimize data transfer between CPU and GPU when possible

Better not to send back from GPU to CPU

Be Aware of Cache Efficiency

- Take full advantage of hardware parallel computing
- Try to avoid the von Neumann bottleneck



GPU Bounds and Performance

Application performance is limited by:

- **Memory Bounds**
- **ALU Bounds**
- **TMU (Texture Mapping Unit) Bound**
- **BW (Bandwidth) Bound**

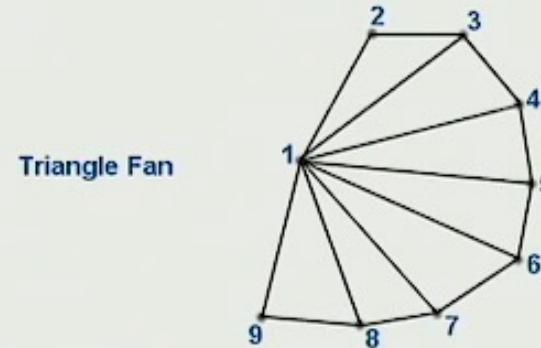
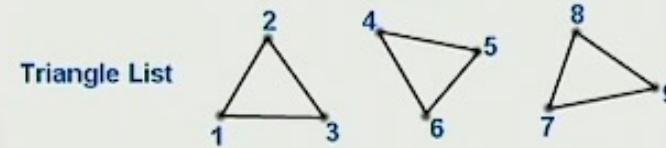
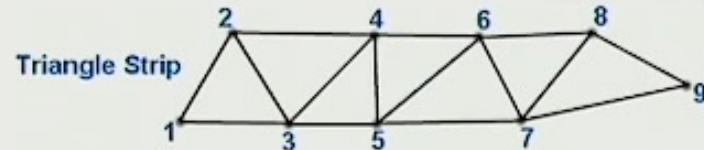
Vertex and Index Buffer

- **Vertex Data**

- Vertex declaration
- Vertex buffer

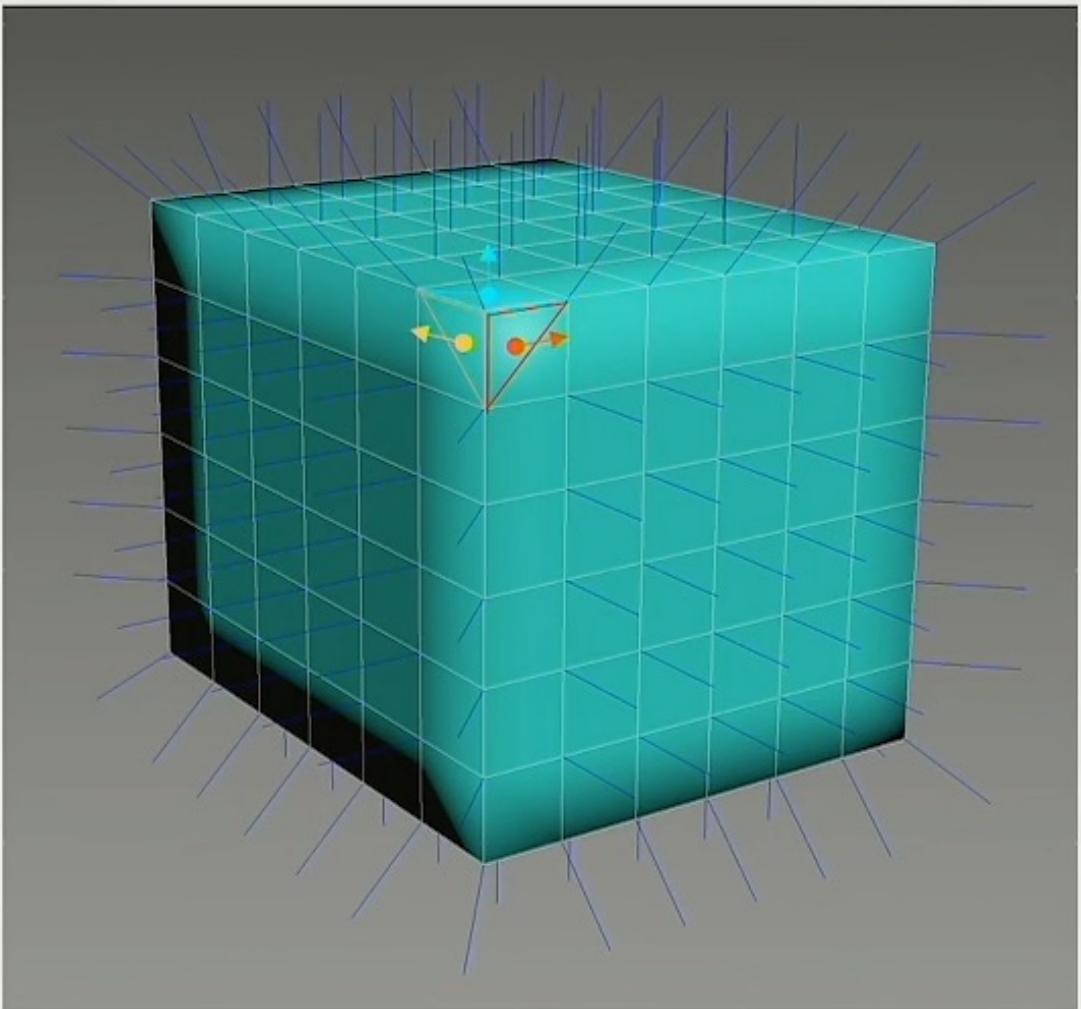
- **Index Data**

- Index declaration
- Index buffer

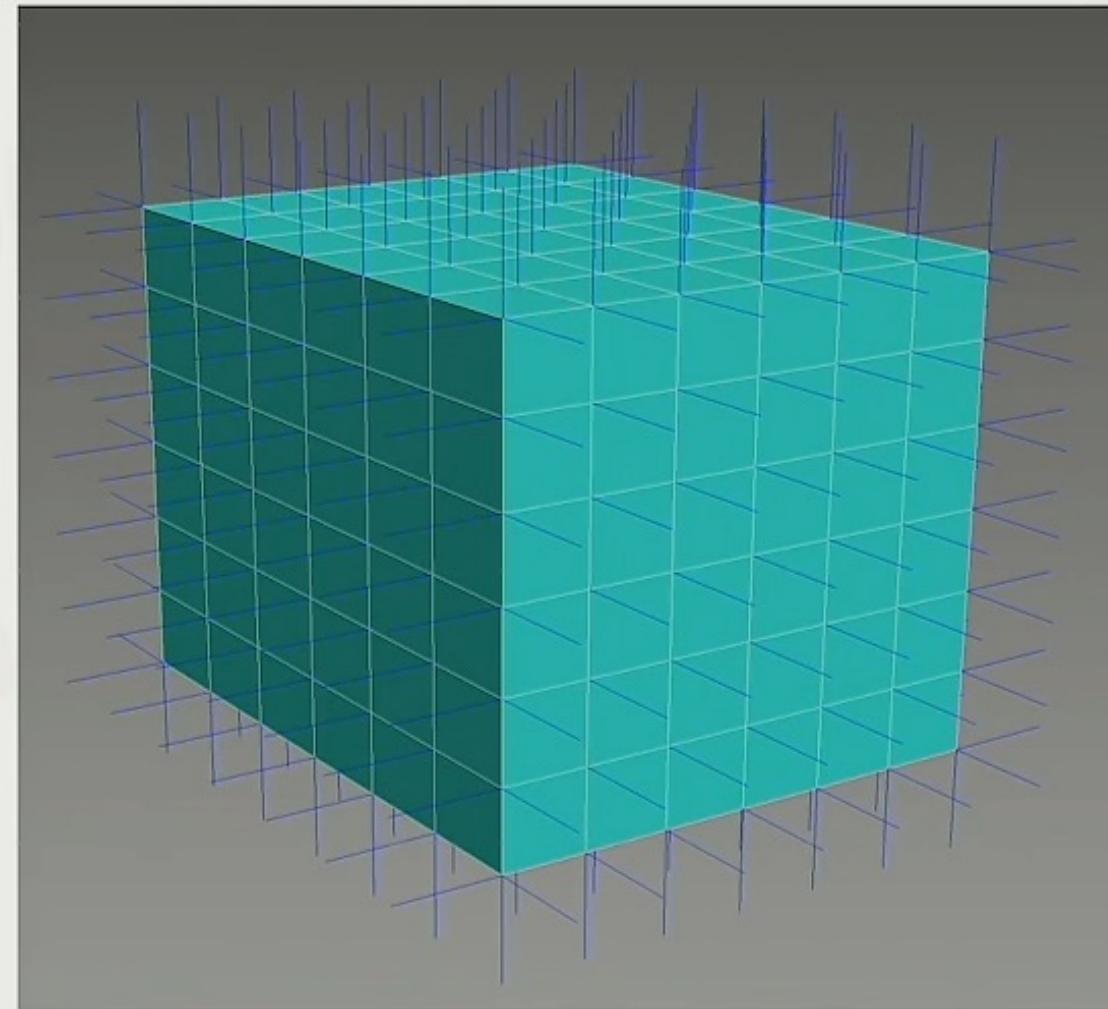


PrimitiveType	TriangleList							
ElementStartAndCount			0		6			
IndexBuffer	0	1	2	2	3	0		
VertexPosition	(x0,y0,z0)	(x1,y1,z1)	(x2,y2,z2)	(x3,y3,z3)				
VertexUV	(u0,v0)	(u1,v1)	(u2,v2)	(u3,v3)				
VertexNormal	(nx0,ny0,nz0)	(nx1,ny1,nz1)	(nx2,ny2,nz2)	(nx3,ny3,nz3)				

Why We Need Per-Vertex Normal



Interpolate vertex normal by triangle normal

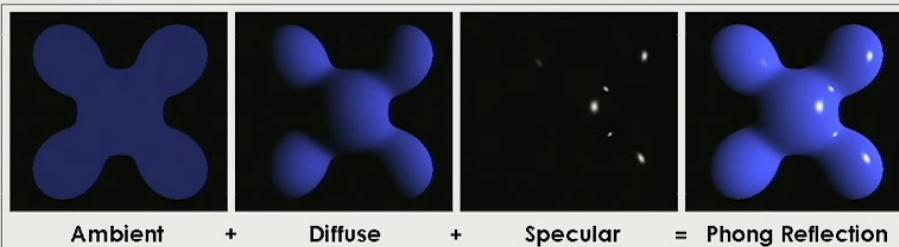


Per-Vertex normals necessary

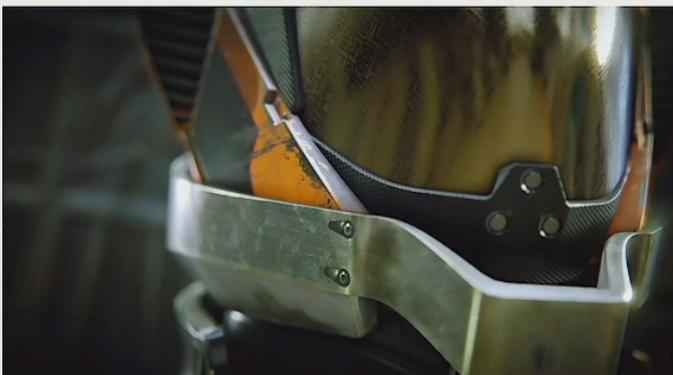
Same vertex different normal for sharp angle.

Materials

Famous Material Models



Phong Model



PBR Model - Physically based rendering



Subsurface Material - Burley SubSurface Profile



Base

Smooth metal

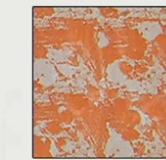
Glossy paint

Rough stone

Transparent glass

Determine the appearance of objects, and how objects interact with light

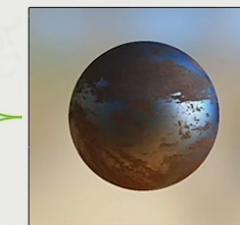
Textures in Materials



ALBEDO



NORMAL



METALLIC



ROUGHNESS

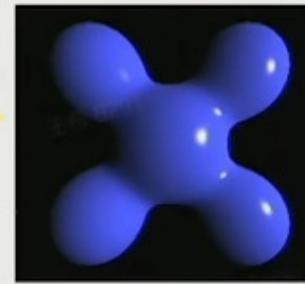


AO

Variety of Shaders

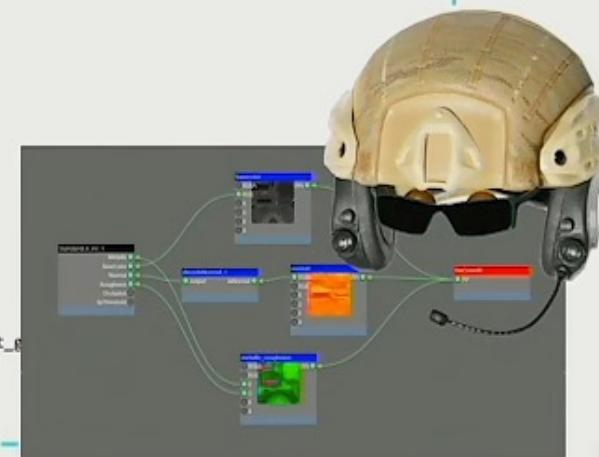
Fix Function Shading Shaders

```
float4 PSMain(PixelInput input) : SV_TARGET
{
    float3 world_normal = normalize(input.world_normal);
    float3 world_view_dir = normalize(world_space_camera_pos - input.world_pos);
    float3 world_light_reflection_dir = normalize(reflect(-world_light_dir, world_normal));
    float3 ambient = ambient_color * material.ambient;
    float3 diffuse = max(0, dot(world_normal, world_light_dir)) *
        diffuse_color * material.diffuse;
    float3 specular = pow(max(0, dot(world_light_reflection_dir, world_view_dir)), shininess) *
        specular_color * material.specular;
    float3 emissive = material.emissive;
    return float4(ambient + diffuse + specular + emissive, 1.0f);
}
```



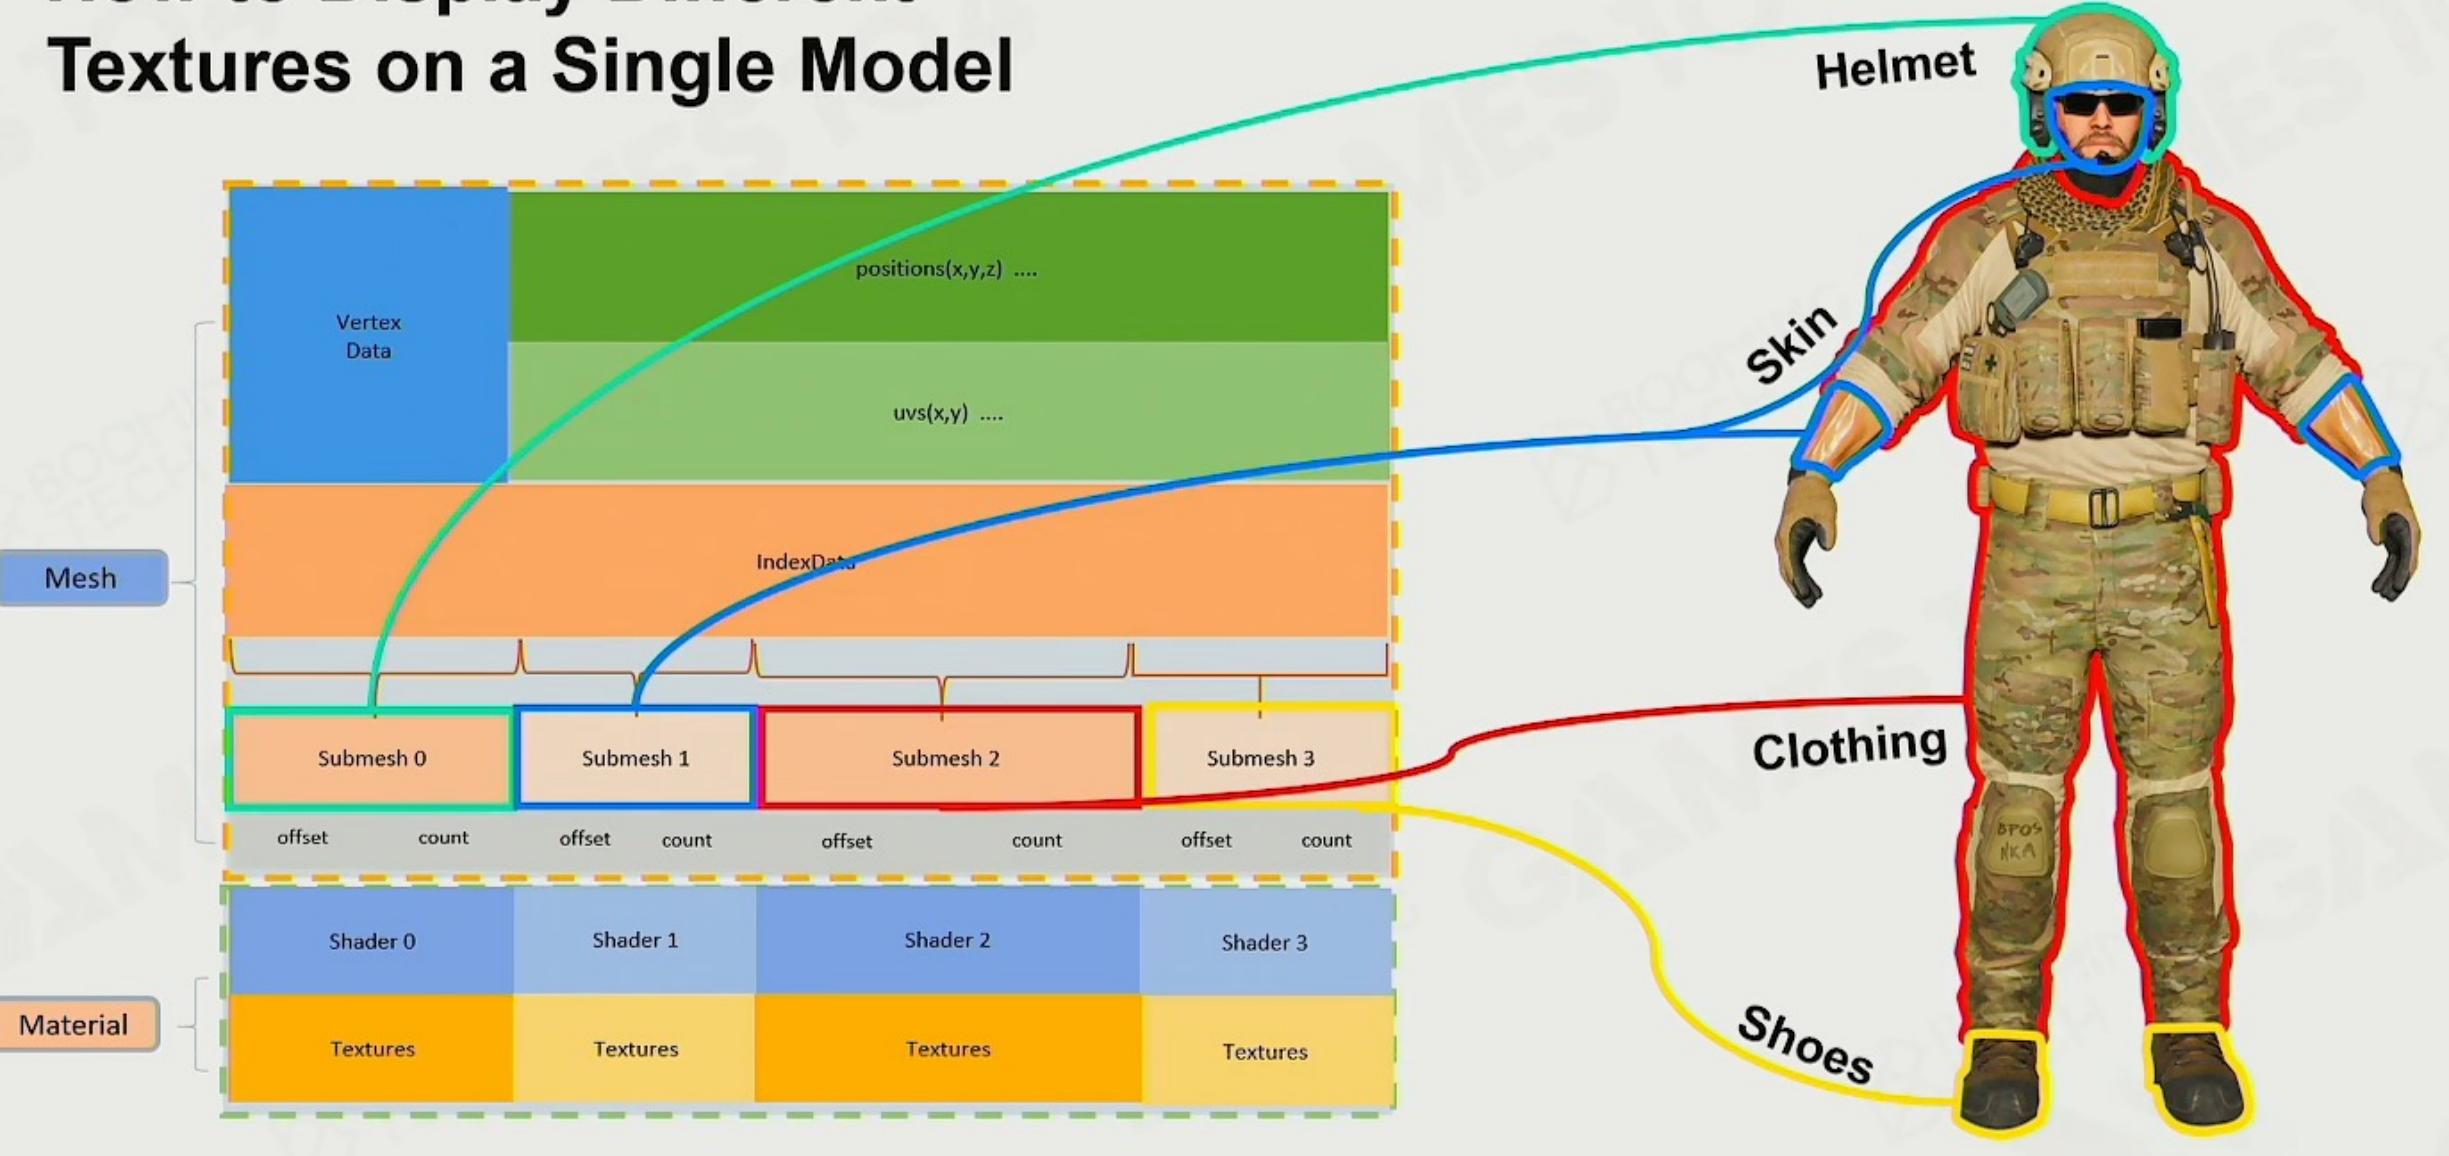
```
104     PixelOutGbuffer PS_Entry_deferred(PixelInput input)
105     {
106         float3 T = input.world_tangent.xyz;
107         float3 N = normalize(input.world_geo_normal);
108         float3 B = cross(input.world_geo_normal, input.world_tangent.xyz)
109             * input.world_tangent.w;
110
111         T -= dot(T, N) * N;
112         T = normalize(T);
113         B -= dot(B, N) * N + dot(B, T) * T;
114         B = normalize(B);
115
116         float3x3 TBN; TBN[0] = T; TBN[1] = B; TBN[2] = N;
117
118         GBufferData gbuffer_data;
119         initializeGBufferData(gbuffer_data);
120
121         //albedo
122         float4 albedo_opacity_value = CHAOS_SAMPLE_TEX2D(albedo_opacity_map, input.tex0);
123         gbuffer_data.albedo = albedo_opacity_value.rgb;
124
125         //normal
126         float3 normal_value = decodeNormalFromNormalMapValue(normal_map.rgb).rgb;
127         gbuffer_data.world_normal = normalize(mul(normal_value.rgb, TBN));
128
129         //specular
130         float4 specular_glossiness_value = CHAOS_SAMPLE_TEX2D(specular_glossiness_map, input.tex0);
131         gbuffer_data.reflectance = specular_glossiness_value.rgb;
132
133         //smoothness
134         gbuffer_data.smoothness = specular_glossiness_value.r;
135
136         //ao
137         gbuffer_data.ao = occlusion;
138
139         //opacity
140         float albedo_opacity_value = albedo_opacity_value.a;
141
142         float alpha_clip_value = alpha_clip;
143
144         clip(albedo_opacity_value - alpha_clip_value);
145
146         PixelOutGbuffer out_gbuffer = (PixelOutGbuffer)0;
147         EncodeGBuffer(gbuffer_data, out_gbuffer.GBufferA, out_gbuffer.GBufferB, out_gbuffer.GBufferC);
148
149     }
150 }
```

Custom Shaders



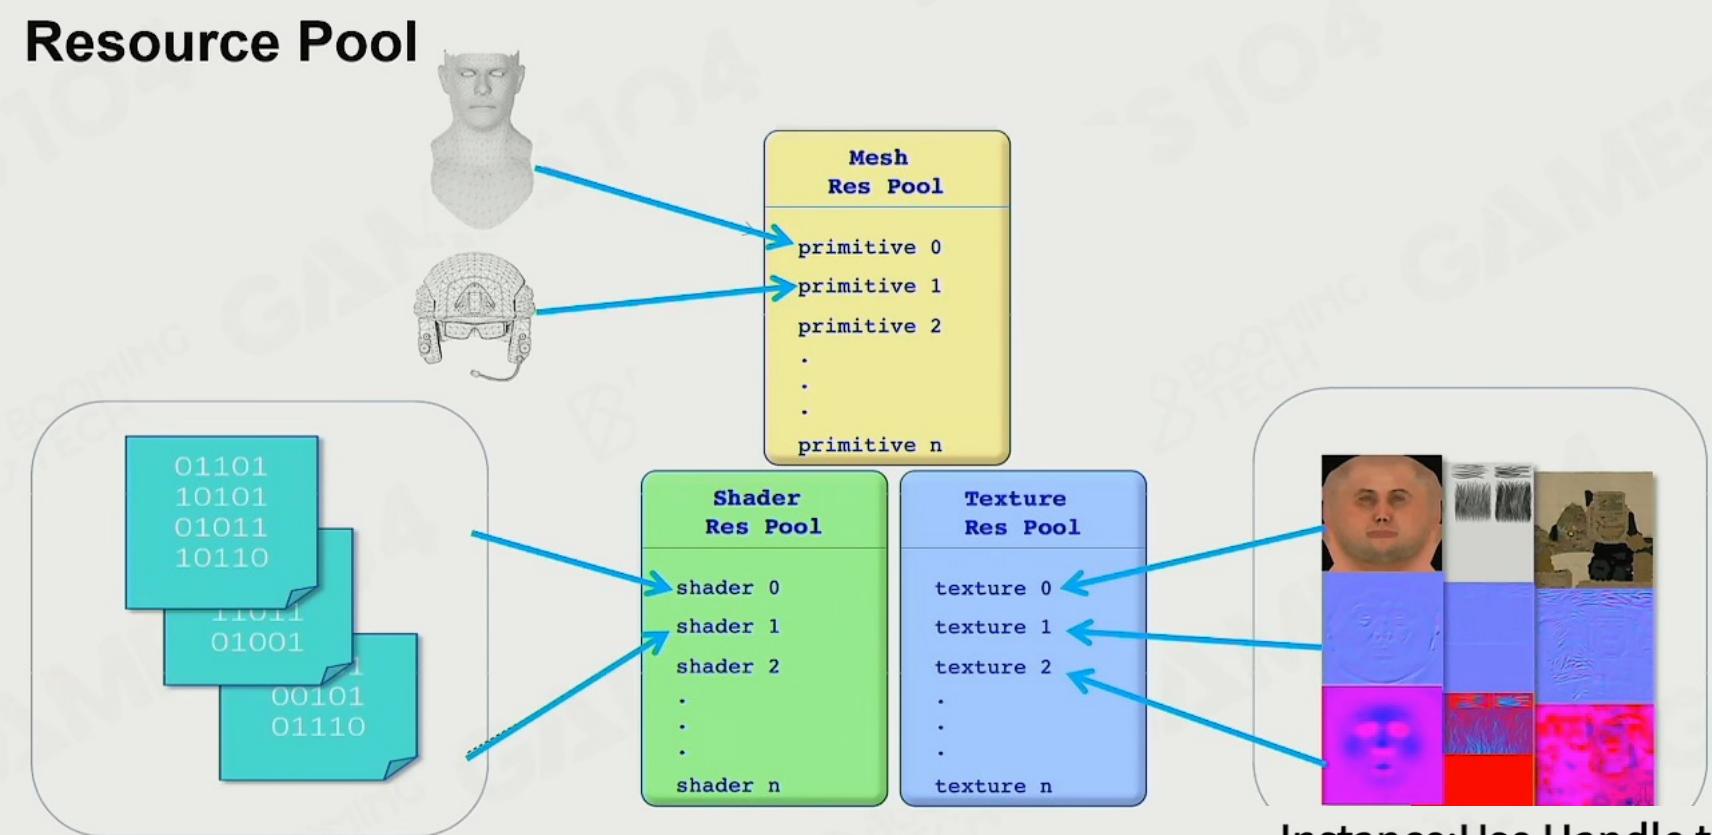
Shader code: source and also data.

How to Display Different Textures on a Single Model

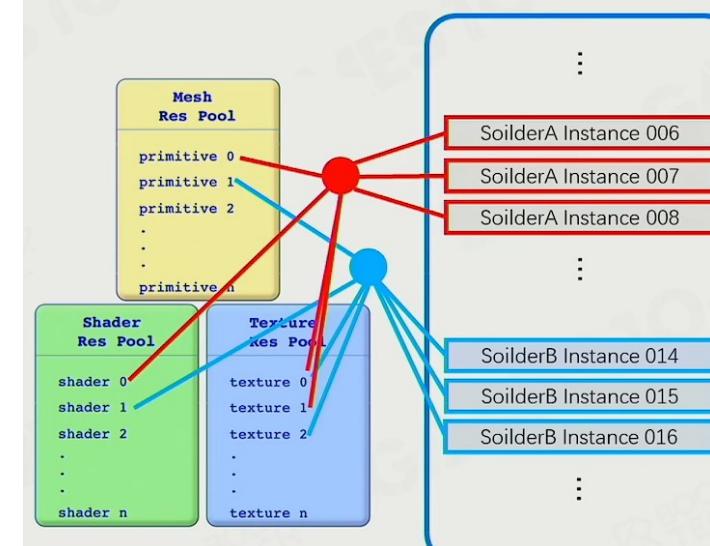


submesh

Resource Pool



Instance: Use Handle to Reuse Resources



Sort by Material

```
Initialize Resource Pools  
Load Resources  
  
Sort all Submeshes by Materials  
  
for each Materials  
    Update Parameters  
    Update Textures  
    Update Shader  
    Update VertexBuffer  
    Update IndexBuffer  
    for each Submeshes  
        Draw Primitive  
    end  
end
```



GPU Batch Rendering



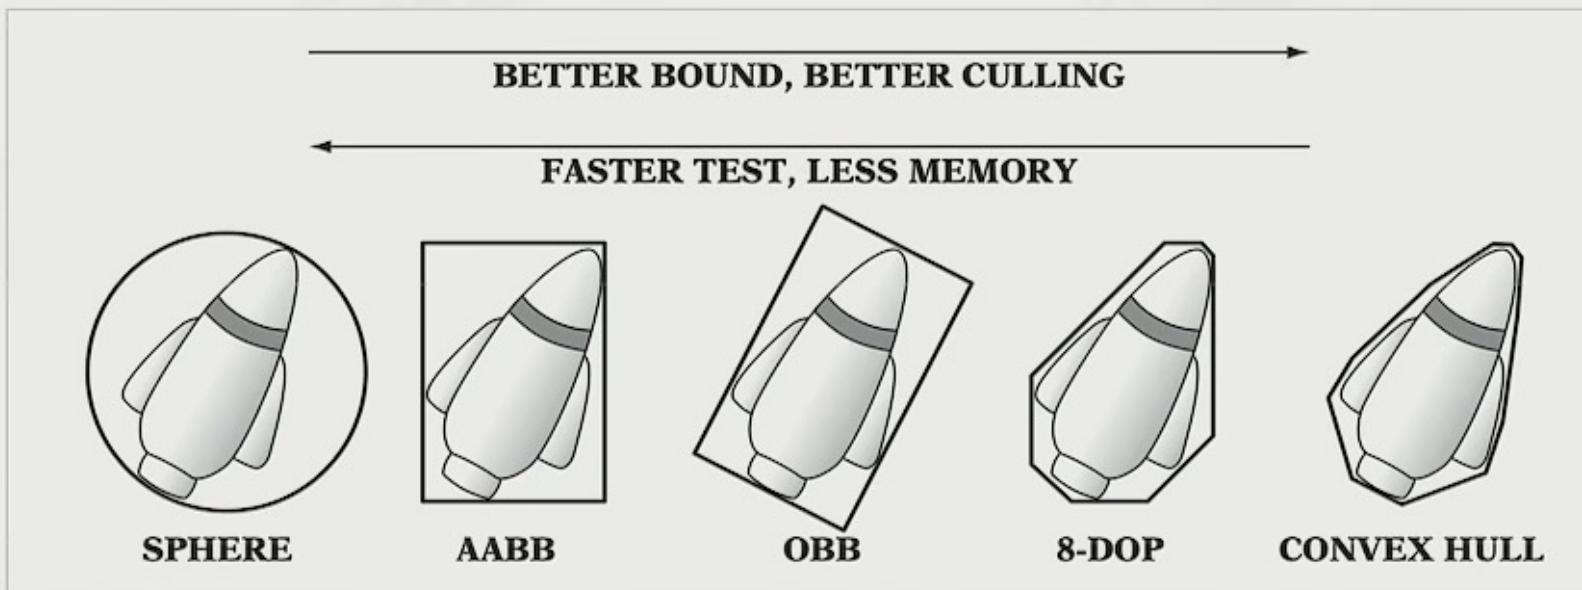
```
struct batchData  
{  
    SubmeshHandle m_submesh_handle;  
    MaterialHandle m_material_handle;  
    std::vector<PerInstanceData> m_per_instance_data;  
    unsigned int m_instance_count;  
};  
  
Initialize Resource Pools  
Load Resources  
  
Collect batchData with same submesh and material  
  
for each BatchData  
    Update Parameters  
    Update Textures  
    Update Shader  
    Update VertexBuffer  
    Update IndexBuffer  
    Draw Instance  
end
```

Q :

What if group rendering all instances with identical submeshes and materials together?

Visibility Culling 可见性裁剪

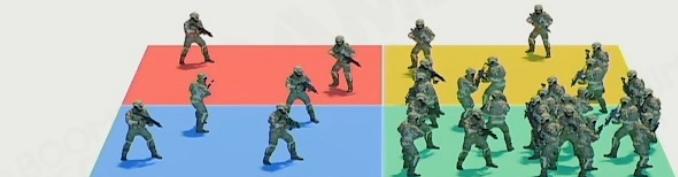
Using the Simplest Bound to Create Culling



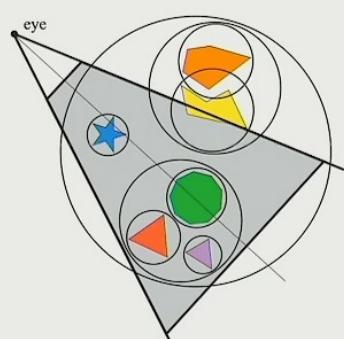
- Inexpensive intersection tests
- Tight fitting
- Inexpensive to compute
- Easy to rotate and transform
- Use little memory

Hierarchical View Frustum Culling

BVH



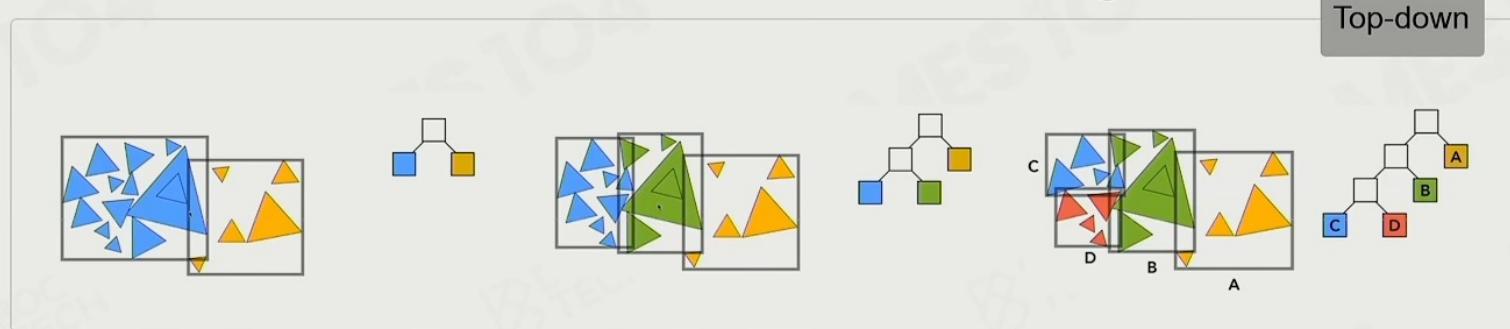
Quad Tree Culling



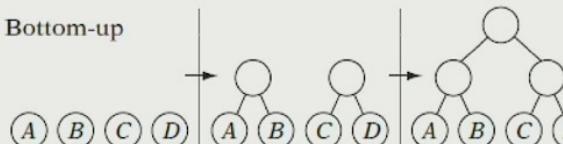
BVH (Bounding Volume Hierarchy) Culling

Construction and Insertion of BVH in Game Engine

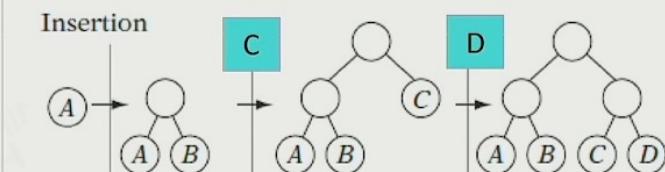
Top-down



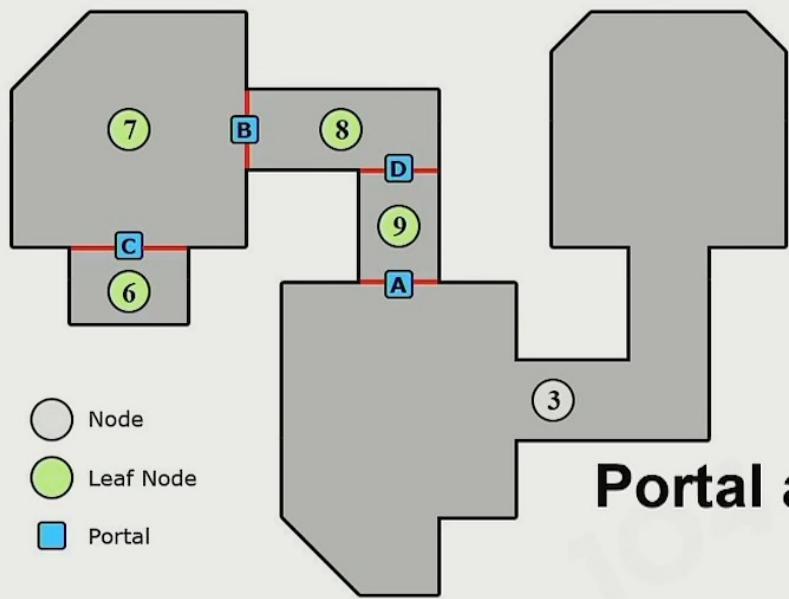
Bottom-up



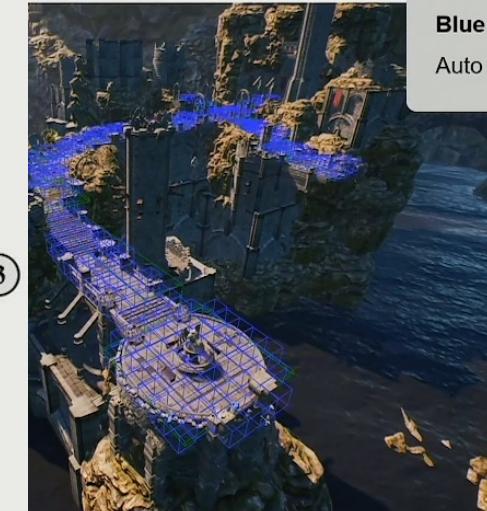
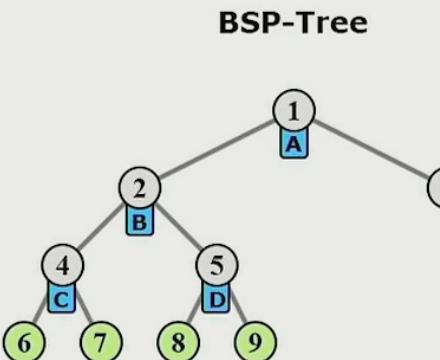
Incremental tree-insertion



PVS (Potential Visibility Set)



Idea used in Resource loading



Green box:

The area to determine the potential visibility where you need.

Blue cells:

Auto generated smaller regions of each green box.

```
for each GreenBoxs  
  for each BlueCells  
    do ray casting between box and cell  
  end  
end
```

Pros

- Much faster than BSP / Octree
- More flexible and compatible
- Preload resources by PVS

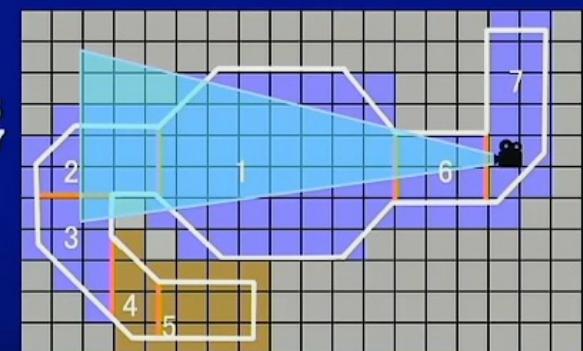
```
for each portals  
  getSamplingPoints();  
  for each portal faces  
    for each leaf  
      do ray casting between portal face and leaf  
    end  
  end  
end
```

Generate PVS data from portal:

Determine potentially visible leaf nodes immediately from portal

Potentially Visible Set

7: 6 1 2 3
6: 1 7 2 3
1: 6 2 7 3
2: 1 3 6 7 4 5
3: 4 2 5 1 6 7
5: 4 3 2
4: 3 5 2



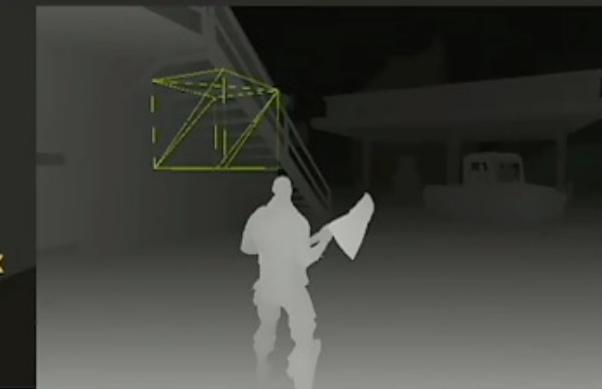
GPU Culling

Render base pass

disable color write

disable depth write

enable depth test



select occludees

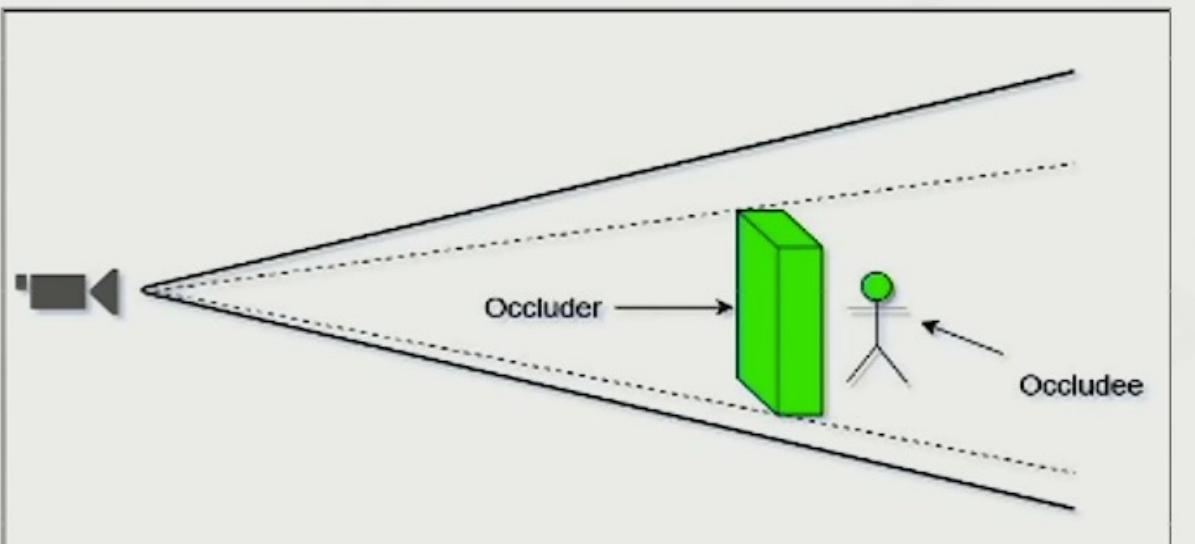
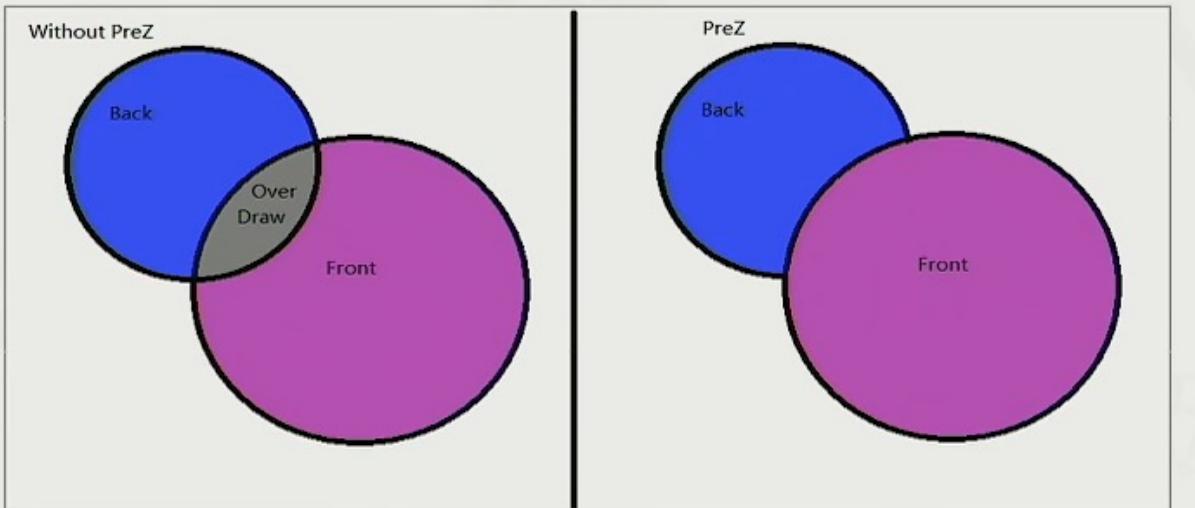
batch occludees

for each occludee

Begin Query

Render occludee bbox

End Query



Seems like Z-buffer

Texture Compression

- **Traditional image compression like JPG and PNG**

- Good compression rates
- Image quality
- Designed to compress or decompress an entire image

- **In game texture compression**

- Decoding speed
- Random access
- Compression rate and visual quality
- Encoding speed

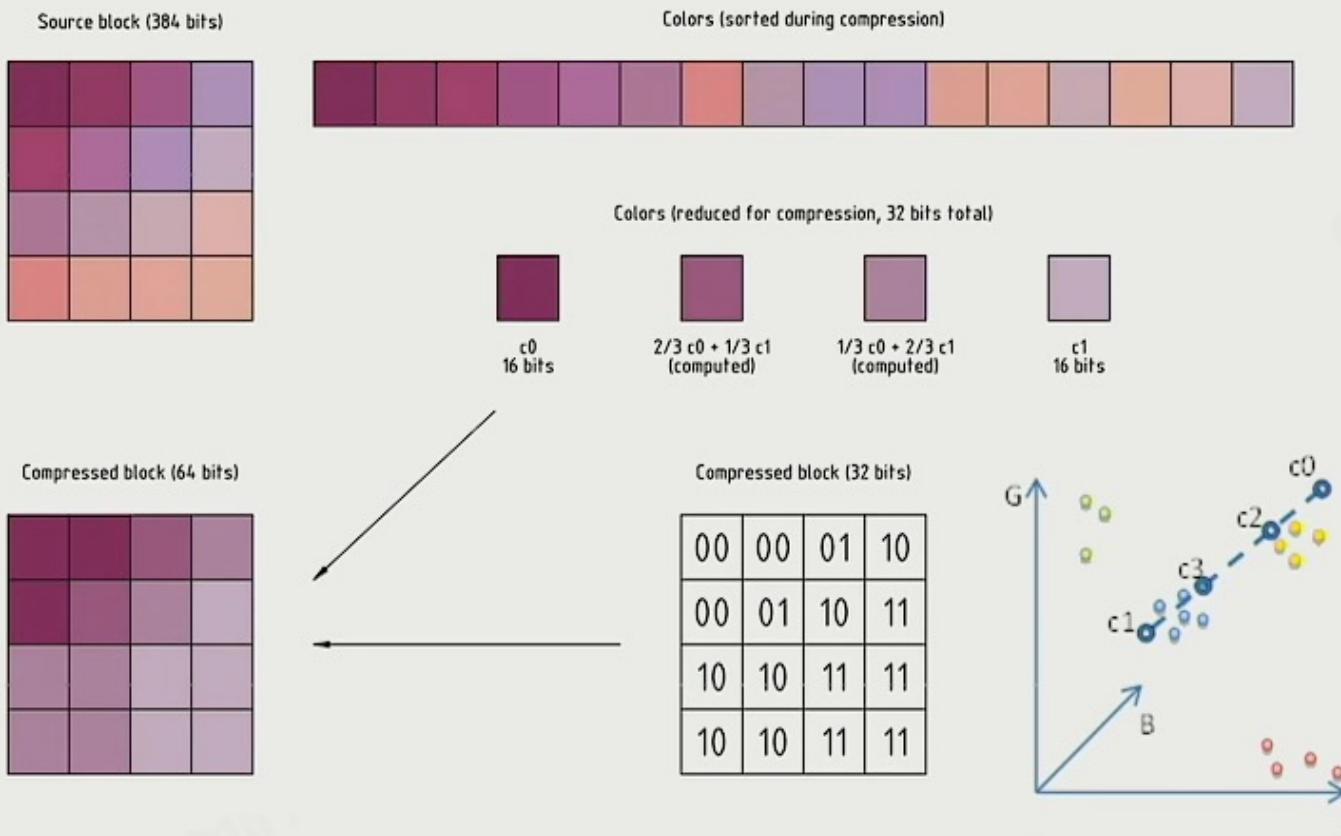


Sample JPEG format texture

Block Compression

Common block-based compression format

- On PC, BC7 (modern) or DXTC (old) formats
- On mobile, ASTC (modern) or ETC / PVRTC (old) formats



Modeling

Comparison of Authoring Methods

	Polymodeling	Sculpting	Scanning	Procedural modeling
Sample				
Advantage	Flexible	Creative	Realistic	Intelligent
Disadvantage	Heavy workload	Large volume of data	Large volume of data	Hard to achieve

Maya, max, blender

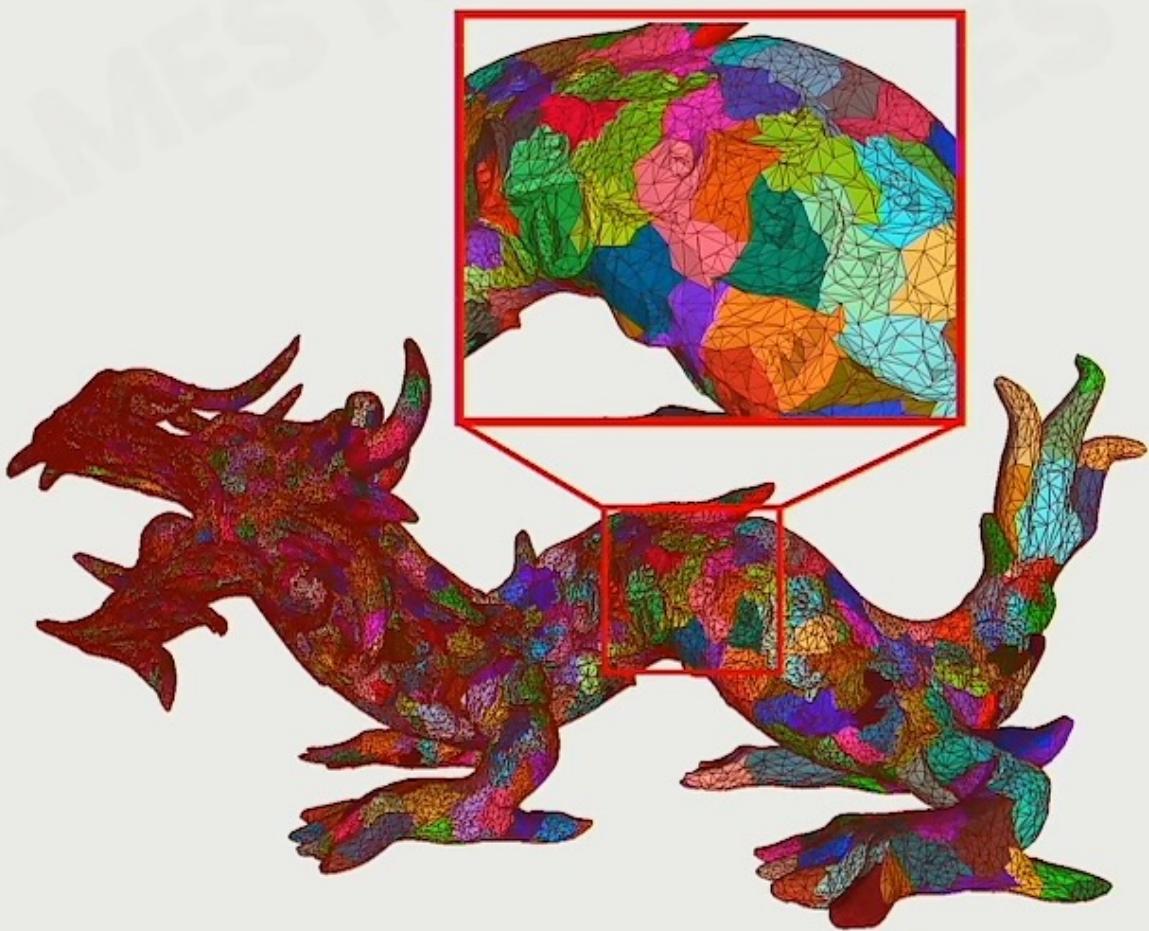
zbrush

Houdini, unreal

Cluster-Based Mesh Pipeline

GPU-Driven Rendering Pipeline (2015)

- Mesh Cluster Rendering
 - Arbitrary number of meshes in single drawcall
 - GPU-culled by cluster bounds
 - Cluster depth sorting

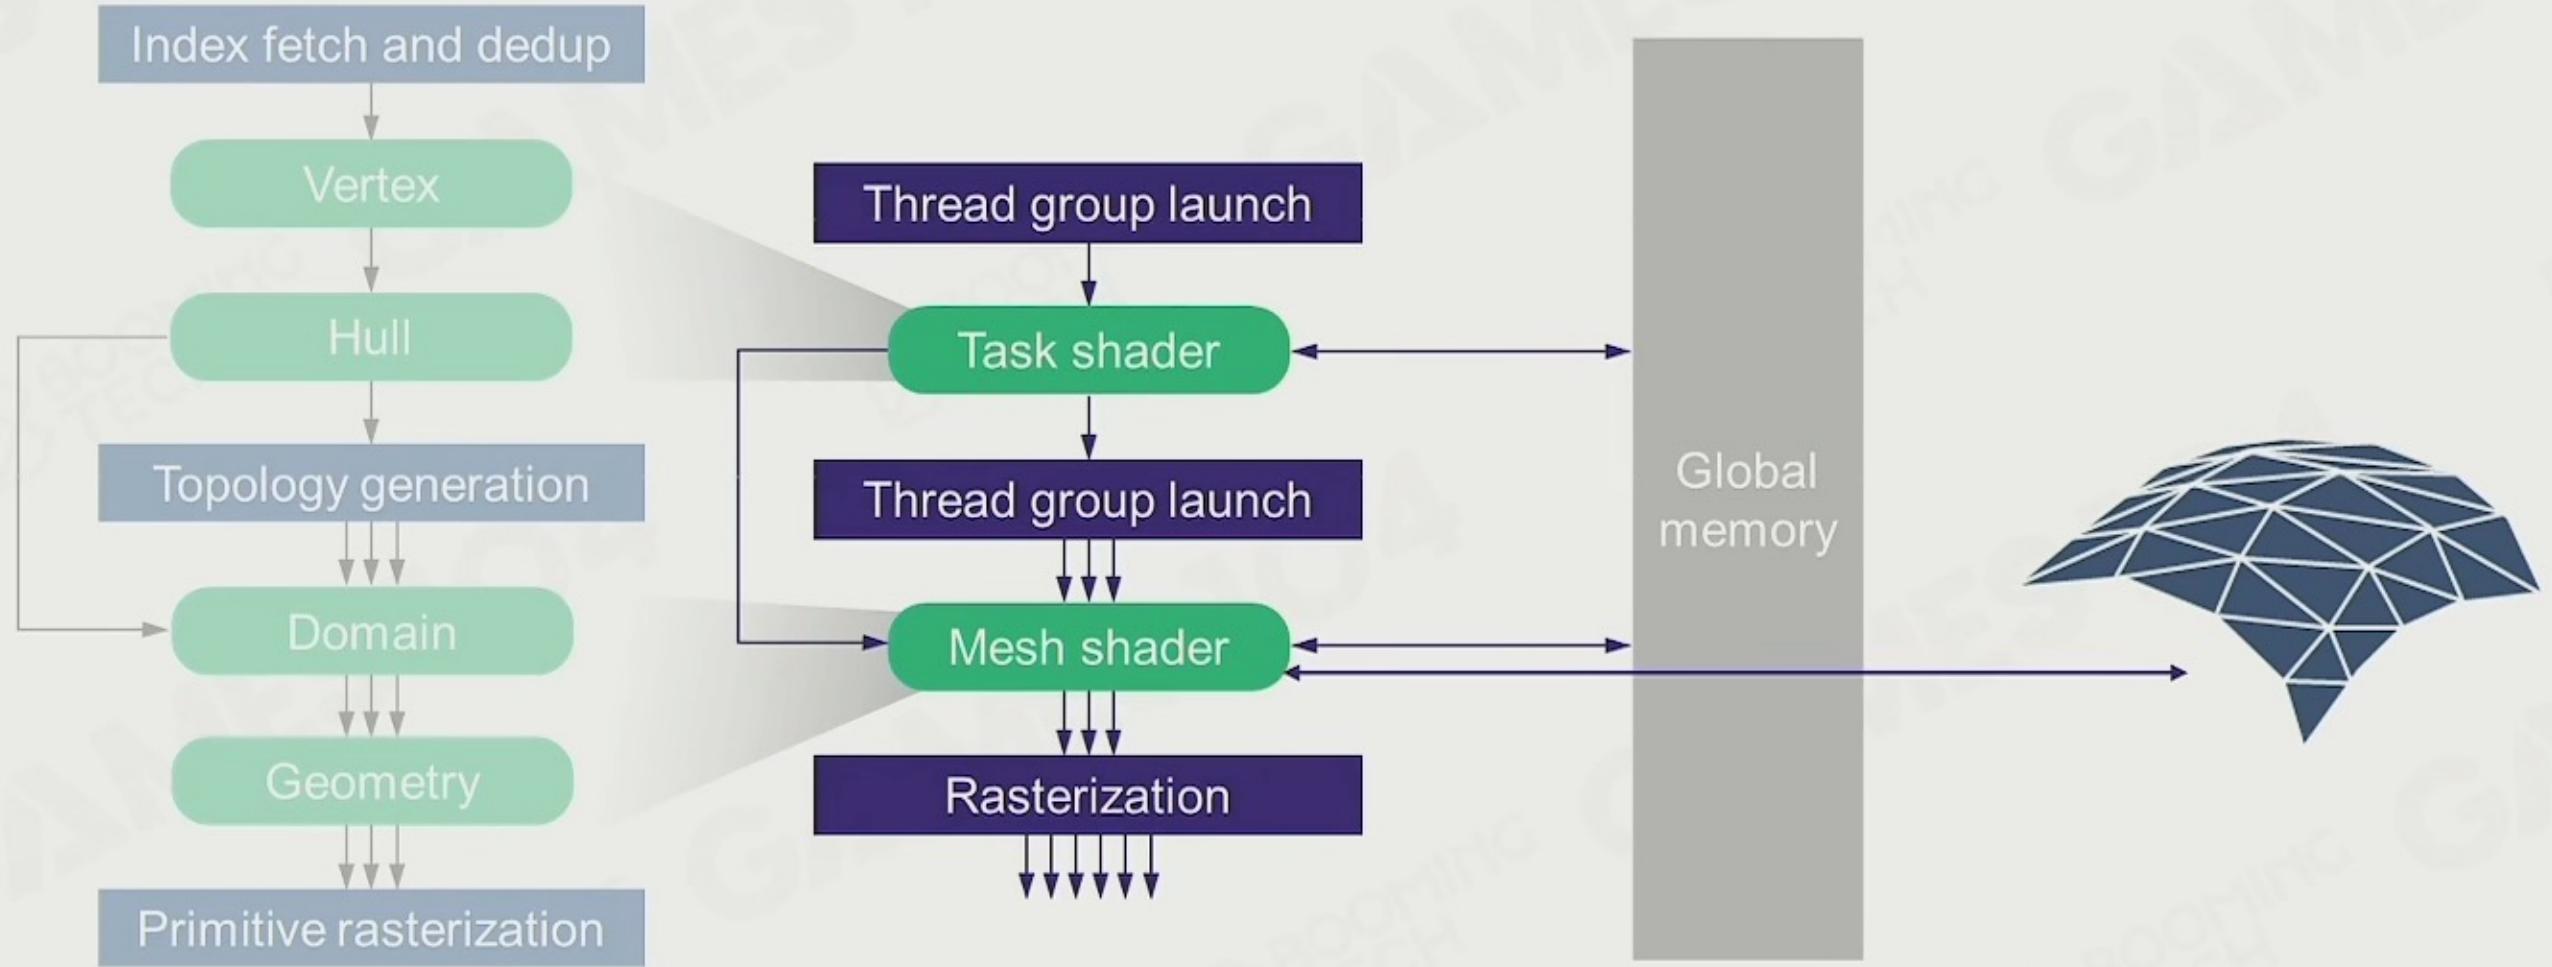


Geometry Rendering Pipeline Architecture (2021)

- Rendering primitives are divided as:
 - Batch: a single API draw (`drawIndirect` / `drawIndexIndirect`), composed of many Surfs
 - Surf: submeshes based on materials, composed of many Clusters
 - Cluster: 64 triangles strip

New trend

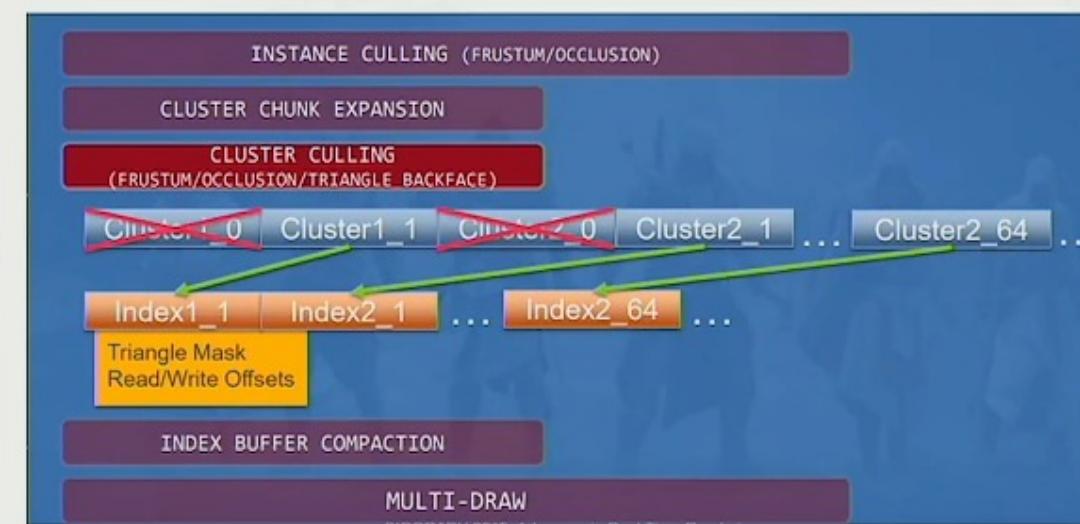
Programmable Mesh Pipeline



GPU Culling in Clustered-Based Mesh



350k triangles to 2791 clusters



GPU Pipeline

Nanite

- Hierarchical LOD clusters with seamless boundary
- Don't need hardware support, but using a hierarchical cluster culling on the precomputed BVH tree by persistent threads (CS) on GPU instead of task shader



Unreal e.g.

(Lighting, Materials, and Shaders in rendering

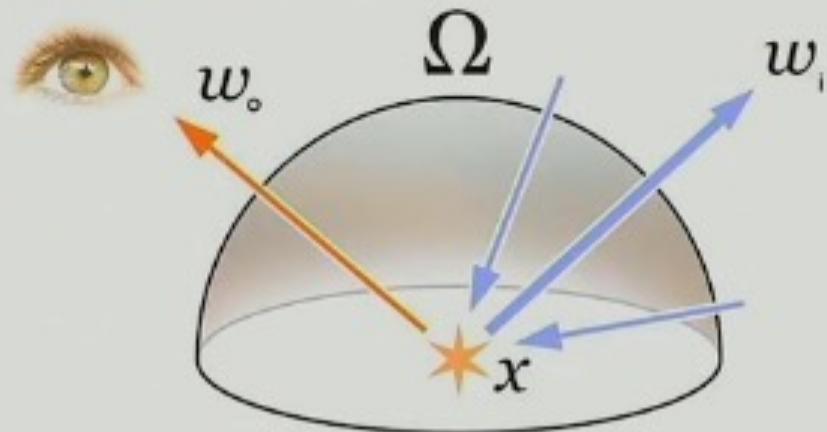
Summary of Popular AAA Rendering

- Lightmap + Lightprobe
- PBR + IBL
- Cascade shadow + VSSM



Lighting

The Rendering Equation

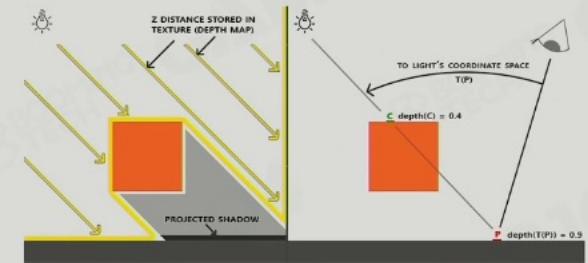


$$L_o(x, \omega_o) = L_e(x, \omega_o) + \int_{H^2} f_r(x, \omega_o, \omega_i) L_i(x, \omega_i) \cos \theta_i d\omega_i$$

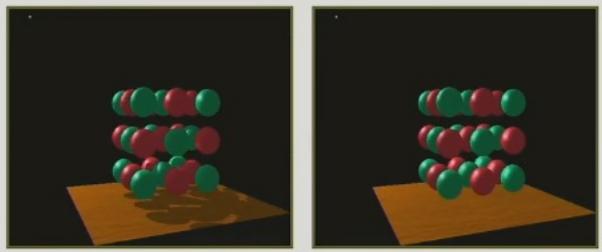
Annotations for the rendering equation:

- outgoing/observed radiance
- emitted radiance (e.g., light source)
- point of interest
- direction of interest
- all directions in hemisphere
- scattering function
- incoming radiance
- angle between incoming direction and normal
- incoming direction

The 1st Challenge: 1a Visibility to Lights



Ray Casting Toward Light Source



Shadow on and off

The 1st Challenge: 1b Light Source Complexity

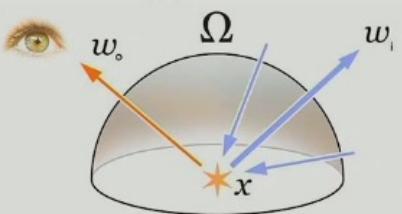


Light Power

面光源，复杂

The 2nd Challenge: How to do Integral Efficiently on Hardware

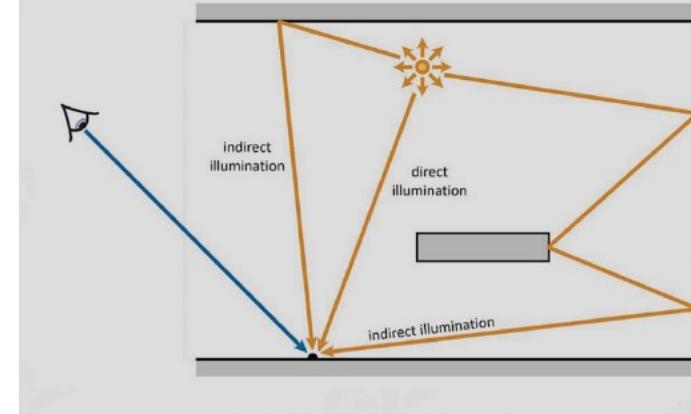
- Brute-force way sampling
- Smarter sampling, i.e., Monte Carlo
- Derive fast analytical solutions
 - Simplify the f_r :
 - Assumptions the optical properties of materials
 - Mathematical representation of materials
 - Simplify the L_i :
 - Deal with directional light, point light and spot light only
 - A mathematical representation of incident light sampling on a hemisphere, for ex: IBL and SH



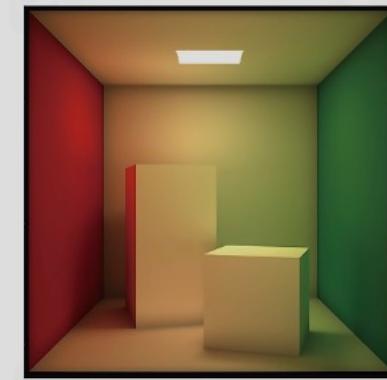
$$L_o(x, \omega_o) = \int_{H^2} f_r(x, \omega_o, \omega_i) L_i(x, \omega_i) \cos \theta_i d\omega_i$$

Expensive integral

The 3rd Challenge: Any matter will be light source



Direct vs. Indirect Illumination



Global Illumination (GI)

recursive

Simple Light Solution

- Using simple light source as main light
 - Directional light in most cases
 - Point and spot light in special case
- Using ambient light to hack others
 - A constant to represent mean of complex hemisphere irradiance
- Supported in graphics API



```
glLightfv(GL_LIGHT0, GL_AMBIENT, light_ambient);
glLightfv(GL_LIGHT0, GL_DIFFUSE, light_diffuse);
glLightfv(GL_LIGHT0, GL_SPECULAR, light_specular);
glLightfv(GL_LIGHT0, GL_POSITION, light_position);
```

Environment Map Reflection

- Using environment map to enhance glossary surface reflection
- Using environment mipmap to represent roughness of surface



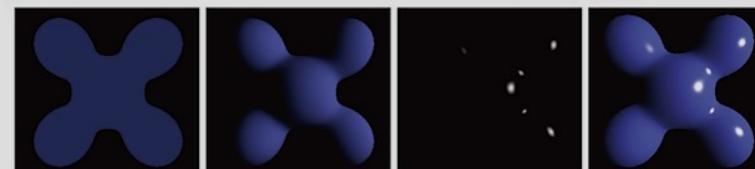
Early stage exploration of image-based lighting

```
void main()
{
    vec3 N = normalize(normal);
    vec3 V = normalize(camera_position - world_position);

    vec3 R = reflect(V, N);

    FragColor = texture(cube_texture, R);
```

Blinn-Phong Materials

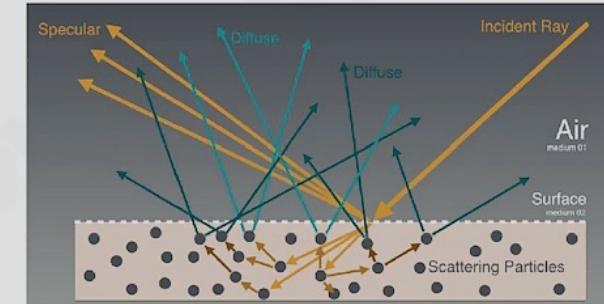


Ambient + Diffuse + Specular = Blinn-Phong Reflection



$$L = L_a + L_d + L_s$$

$$= k_a I_a + k_d (I/r^2) \max(0, \mathbf{n} \cdot \mathbf{l}) + k_s (I/r^2) \max(0, \mathbf{n} \cdot \mathbf{h})^p$$



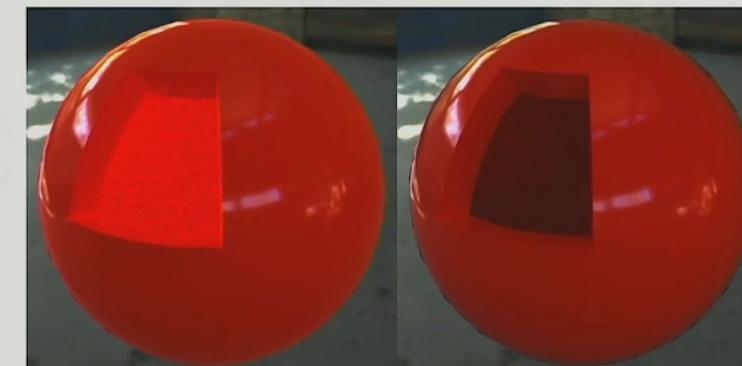
Experience model

```
// set material ambient  
glMaterialfv(GL_FRONT_AND_BACK, GL_AMBIENT, Ka);  
// set material diffuse  
glMaterialfv(GL_FRONT_AND_BACK, GL_DIFFUSE, Kd);  
// set material specular  
glMaterialfv(GL_FRONT_AND_BACK, GL_SPECULAR, Ks);
```

$$L_o(\mathbf{x}, \omega_o) = \int_{H^2} f_r(\mathbf{x}, \omega_o, \omega_i) L_i(\mathbf{x}, \omega_i) \cos \theta_i d\omega_i$$

Problems.

- Not energy conservative
 - Unstable in ray-tracing
- Hard to model complex realistic material



Left non-energy conserving model lead a lot of noise compare Right energy conserving model



Traditional shading

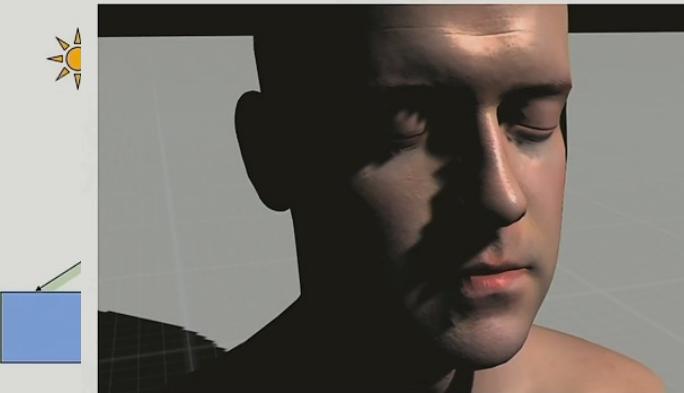
PBR shading

Shadow

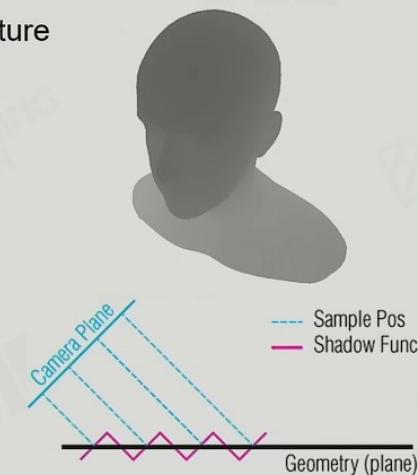
Problem of Shadow Map

- Shadow is nothing but space when the light is blocked by an opaque object
- Already obsolete method
 - planar shadow
 - shadow volume
 - projective texture

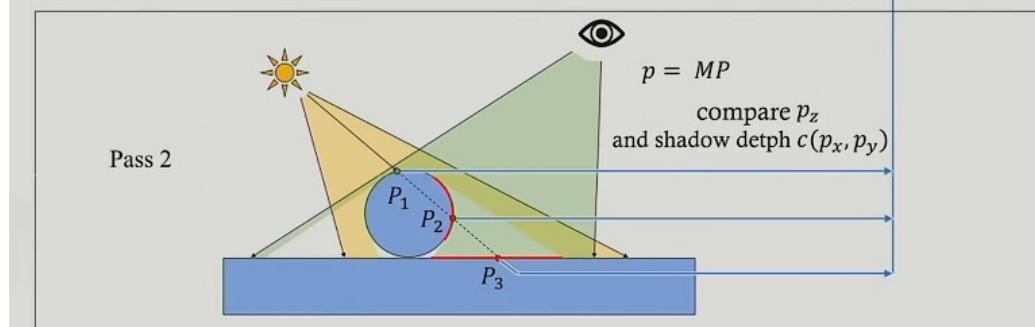
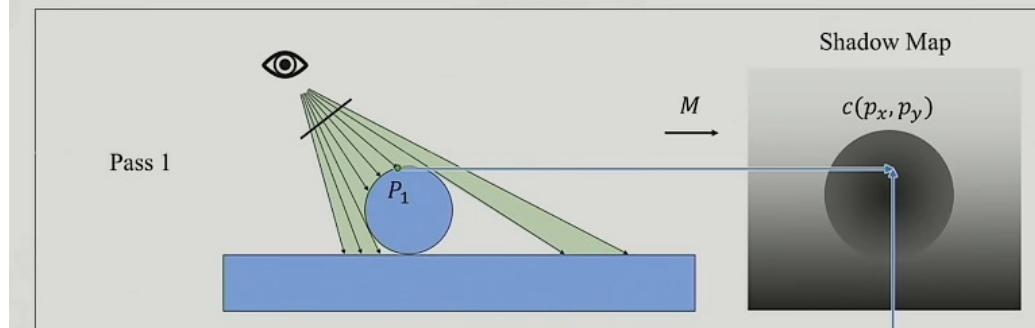
Resolution is limited on texture



Depth precision is limited in texture



Shadow Map



```
//project our 3D position to the shadow map  
vec4 proj_pos = shadow_viewproj * pos;  
  
//from homogeneous space to clip space  
vec2 shadow_uv = proj_pos.xy / proj_pos.w;  
  
//from clip space to uv space  
shadow_uv = shadow_uv * 0.5 + vec2(0.5);  
  
//get point depth (from -1 to 1)  
float real_depth = proj_pos.z / proj_pos.w;  
  
//normalize from [-1..+1] to [0..+1]  
real_depth = real_depth * 0.5 + 0.5;  
  
//read depth from depth buffer in [0..+1]  
float shadow_depth = texture(shadowmap, shadow_uv).x;  
  
//compute final shadow factor by comparing  
float shadow_factor = 1.0;  
if (shadow_depth < real_depth)  
    shadow_factor = 0.0;
```

Basic Shading Solution

- Simple light + Ambient
 - dominant light solves No. 1b
 - ambient and EnvMap solves No. 3 challenges
- Blinn-Phong material
 - solve No. 2 challenge
- Shadow map
 - solve No.1a challenge



Doom3

$$L_r(\mathbf{x}, \vec{\omega}_r) = \int_{H^2} f_r(\mathbf{x}, \vec{\omega}_i, \vec{\omega}_r) L_i(\mathbf{x}, \vec{\omega}_i) \cos \theta_i d\vec{\omega}_i$$

Pre-computed global illumination

Why Global Illumination is Important



Direct illumination



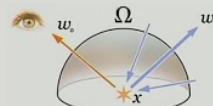
Direct + indirect illumination

How to Represent Indirect Light

- Good compression rate
 - We need to store millions of radiance probes in a level
- Easy to do integration with material function
 - Use polynomial calculation to convolute with material BRDF



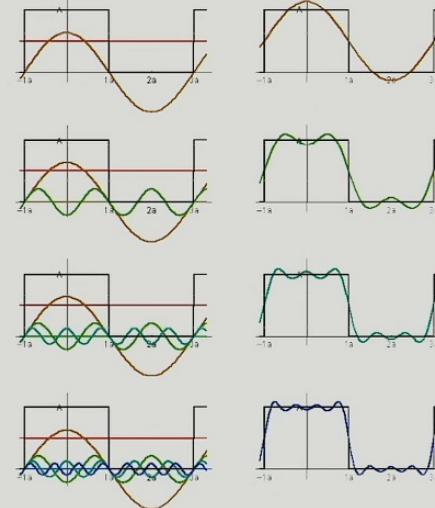
$$L_r(\mathbf{x}, \vec{\omega}_r) = \int_{H^2} f_r(\mathbf{x}, \vec{\omega}_i, \vec{\omega}_r) L_i(\mathbf{x}, \vec{\omega}_i) \cos \theta_i d\vec{\omega}_i$$



Fourier Transform



Joseph Fourier 1768 - 1830



$$f(x) = \frac{A}{2} + \frac{2A \cos(t\omega)}{\pi} - \frac{2A \cos(3t\omega)}{3\pi} + \frac{2A \cos(5t\omega)}{5\pi} - \frac{2A \cos(7t\omega)}{7\pi} + \dots$$

Convolution Theorem



$$\star \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} =$$

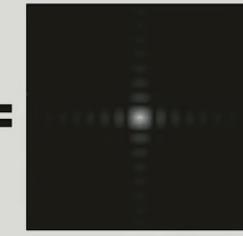


Spatial
Domain

Fourier
Transform

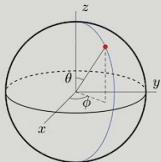
Frequency
Domain

Inv. Fourier
Transform

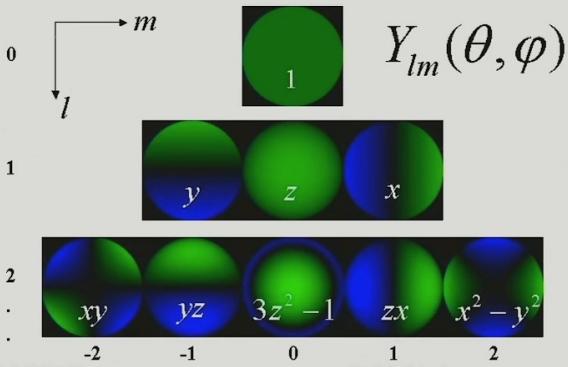


Spherical Harmonics

$$Y_{lm}(\theta, \phi) = N_{lm} P_{lm}(\cos \theta) e^{im\phi}$$

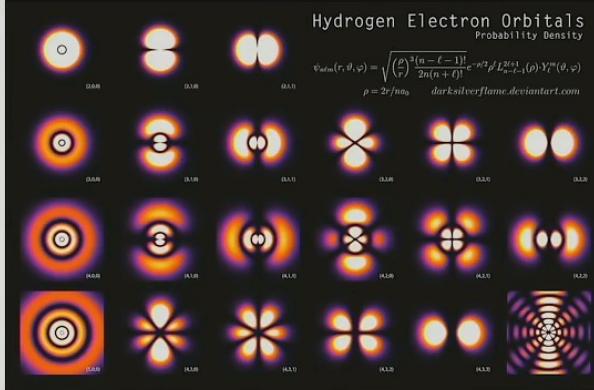
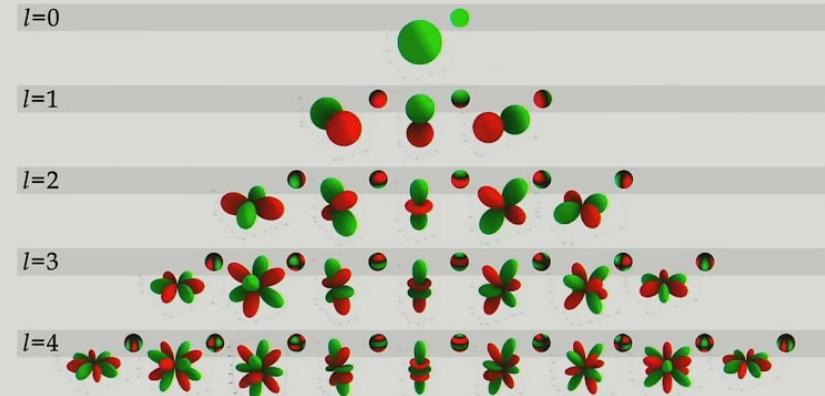


$$\begin{aligned} x &= \sin \theta \cos \phi \\ y &= \sin \theta \sin \phi \\ z &= \cos \theta \end{aligned}$$



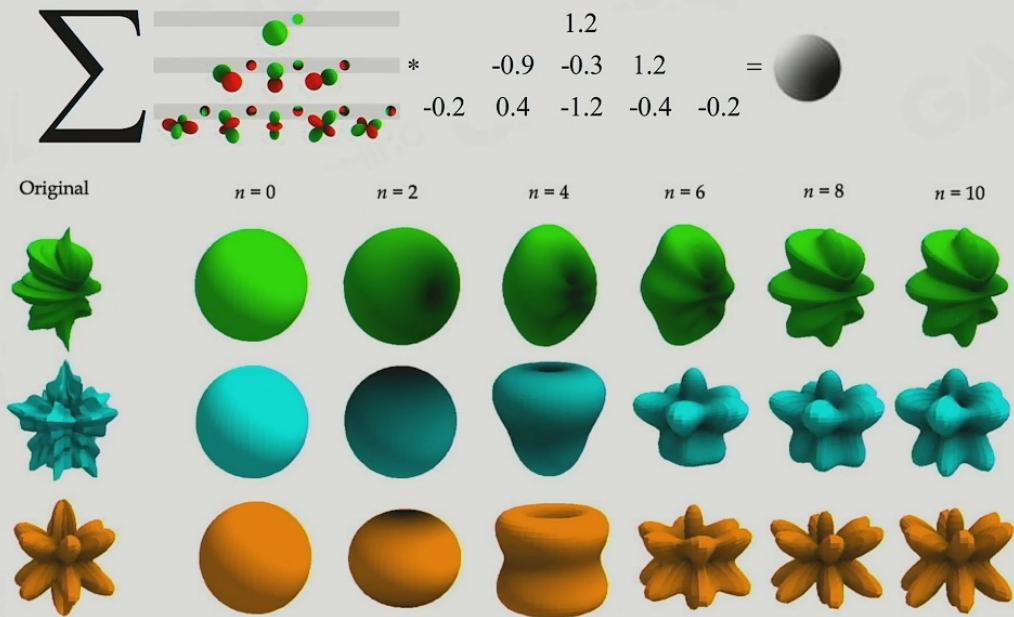
Complex sphere integration can be approximated by quadratic polynomial:

$$\int_{\theta=0}^{\pi} \int_{\phi=0}^{2\pi} L(\theta, \phi) Y_{lm}(\theta, \phi) \sin \theta d\theta d\phi \approx \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}^T M \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$



Spherical Harmonics, a mathematical system analogous to the Fourier transform but defined across the surface of a sphere. The SH functions in general are defined on imaginary numbers

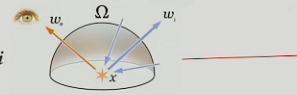
Spherical Harmonics Encoding



Sampling Irradiance Probe Anywhere



$$L_r(\mathbf{x}, \vec{\omega}_r) = \int_{H^2} f_r(\mathbf{x}, \vec{\omega}_i, \vec{\omega}_r) L_i(\mathbf{x}, \vec{\omega}_i) \cos \theta_i d\vec{\omega}_i$$



Store and Shading with SH

- Use 4 RGB textures to store 12 SH coefficients
 - L_0 coefficients in HDR (BC6H texture)
 - L_1 coefficients in LDR (3x BC7 or BC1 textures)
- Total footprint for RGB SH lightmaps:
 - 32 bits (4 bytes) / texel for BC6+BC7, high quality mode
 - 20 bits (2.5 bytes) / texel for BC6+BC1, low quality mode

• `shEvaluateL1` can be merged with `shApplyDiffuseConvolutionL1`

```
SHL1 shEvaluateL1(vec3 p)
{
    float Y0 = 0.282095f; // sqrt(1/fourPi)
    float Y1 = 0.488603f; // sqrt(3/fourPi)
    SHL1 sh;
    sh[0] = Y0;
    sh[1] = Y1 * p.y;
    sh[2] = Y1 * p.z;
    sh[3] = Y1 * p.x;
    return sh;
}
```

```
void shApplyDiffuseConvolutionL1(SHL1 sh)
{
    float A0 = 0.886227f; // pi/sqrt(fourPi)
    float A1 = 1.023326f; // sqrt(pi/3)
    SHL1 sh;
    sh[0] = A0;
    sh[1] *= A1;
    sh[2] *= A1;
    sh[3] *= A1;
}
```

```
SHL1 shEvaluateDiffuseL1(vec3 p)
{
    float AY0 = 0.25f;
    float AY1 = 0.50f;
    SHL1 sh;
    sh[0] = AY0;
    sh[1] = AY1 * p.y;
    sh[2] = AY1 * p.z;
    sh[3] = AY1 + p.x;
    return sh;
}
```

Compress Irradiance Probe to SH1



Source Irradiance Probe

Compressed Irradiance Probe By SH1

Reconstruct Irradiance In Shader

Just RGBA8 color

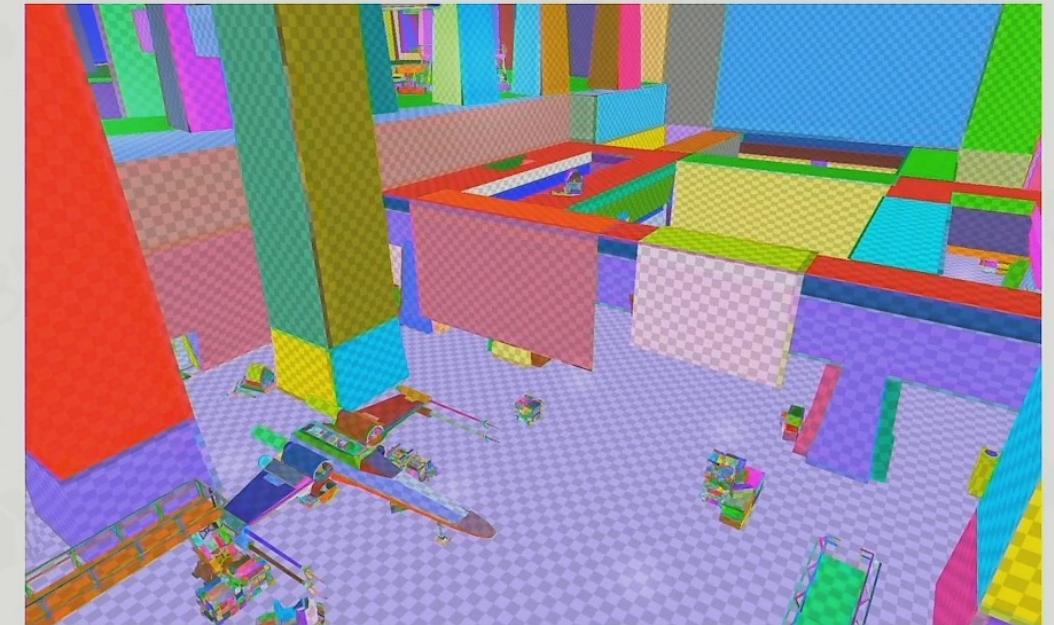
Simple diffuse shading

SH Lightmap : Precomputed GI



- Parameterized all scene into huge 2D lightmap atlas
- Using offline lighting farm to calculate irradiance probes for all surface points
- Compress those irradiance probes into SH coefficients
- Store SH coefficients into 2D atlas lightmap textures

Lightmap: UV Atlas



Lightmap density

- low-poly proxy geometry
- fewer UV charts/islands
- fewer lightmap texels are wasted

Lightmap: Lighting



Indirect lighting, final geometry

- project lightmap from proxies to all LODs
- apply mesh details
- add short-range, high-frequency lighting detail by HBAO

baking

Lightmap: Lighting + Direct Lighting

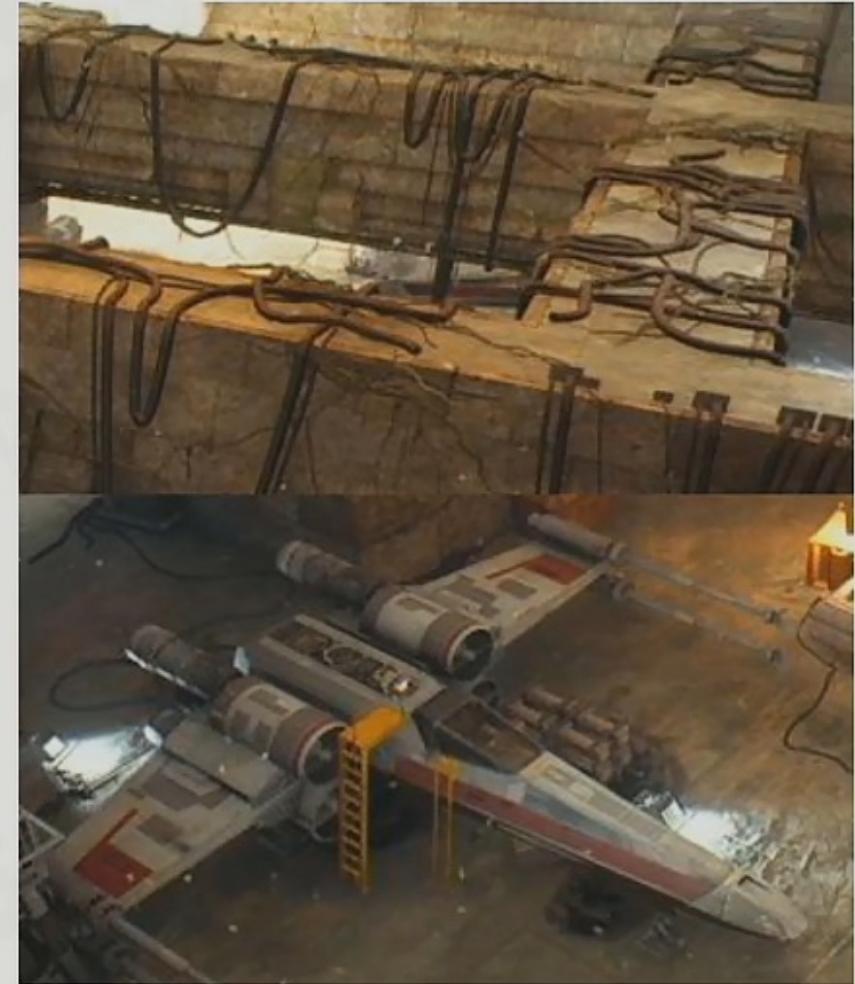


Direct + indirect lighting, final geometry

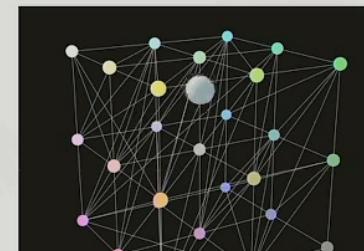
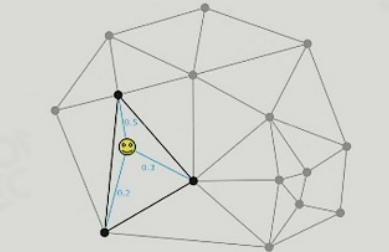
- compute direct lighting dynamically

Lightmap

- **Pros**
 - Very efficient on runtime
 - Bake a lot of fine detail of GI on environment
- **Cons**
 - Long and expensive precomputation (lightmap farm)
 - Only can handle static scene and static light
 - Storage cost on package and GPU

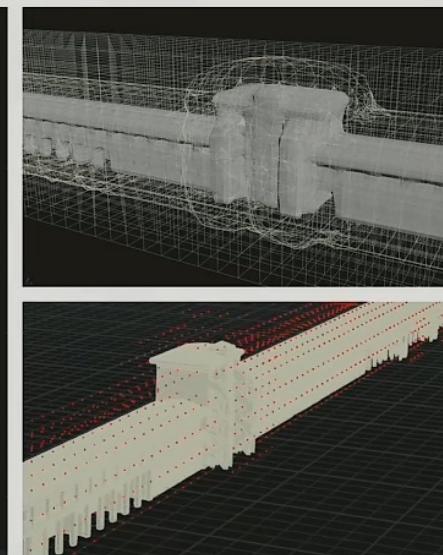
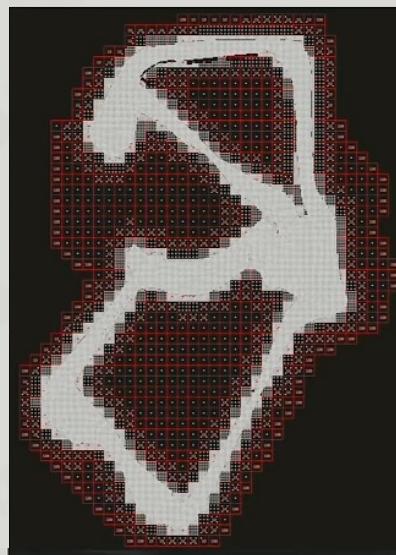


Light Probe: Probes In Game Space



SILVERSTO Light Probe Point Generation

Reflection Probe



Generated by Terrain and Road

Generated by Voxel

Final Light Probe Cloud

Light Probes + Reflection Probes

- **Pros**
 - Very efficient on runtime
 - Can be applied to both static and dynamic objects
 - Handle both diffuse and specular shading
- **Cons**
 - A bunch of SH light probes need some precomputation
 - Can not handle fine detail of GI. I.e, soft shadow on overlapped structures

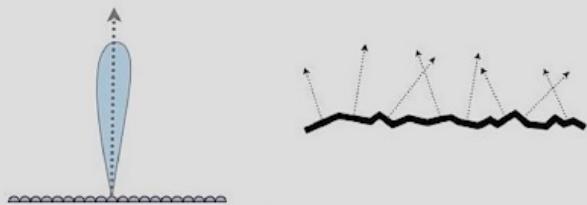
Materials

- Physical – Based

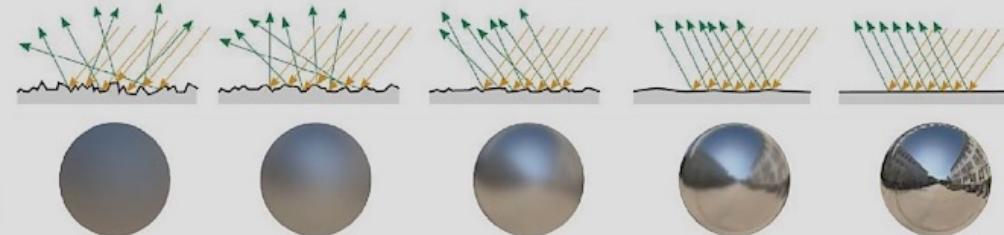
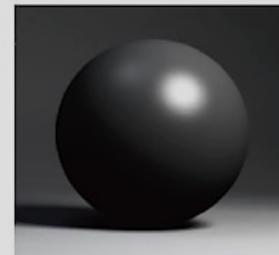
Microfacet Theory

- Key: the **distribution** of microfacets' normals

- Concentrated \iff glossy



- Spread \iff diffuse



BRDF Model Based on Microfacet

$$L_o(x, \omega_o) = \int_{H^2} f_r(x, \omega_o, \omega_i) L_i(x, \omega_i) \cos \theta_i d\omega_i$$



$$f_r = k_d f_{Lambert} + f_{CookTorrance}$$



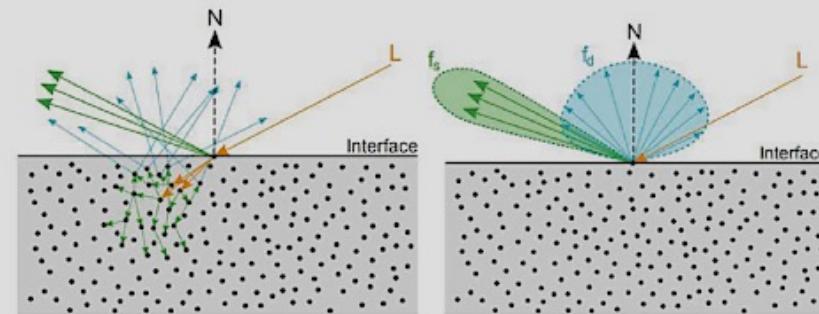
$$f_{Lambert} = \frac{c}{\pi}$$

diffuse

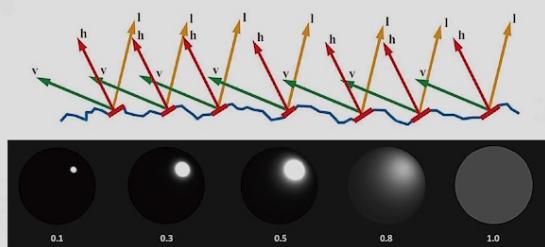
$$f_{CookTorrance} = \frac{DFG}{4(\omega_o \cdot n)(\omega_i \cdot n)}$$

spectral

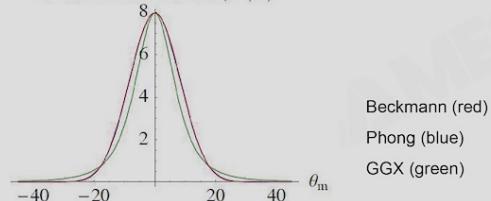
$$L_o(x, \omega_o) = \int_{H^2} \left(k_d \frac{c}{\pi} + \frac{DFG}{4(\omega_o \cdot n)(\omega_i \cdot n)} \right) L_i(x, \omega_i) (\omega_i \cdot n) d\omega_i$$



Normal Distribution Function



Microfacet Distributions, $D(m)$



```
// Normal Distribution Function using GGX Distribution
float D_GGX(float NdotH, float roughness)
{
    float a2 = roughness * roughness;
    float f = (NdotH * NdotH) * (a2 - 1.0) + 1.0;
    return a2 / (PI * f * f);
}
```

$$f_{CookTorrance} = \frac{DFG}{4(\omega_o \cdot n)(\omega_i \cdot n)}$$

$$NDF_{GGX}(n, h, \alpha) = \frac{\alpha^2}{\pi \left((n \cdot h)^2 (\alpha^2 - 1) + 1 \right)^2}$$

Geometric attenuation term (self-shadowing)

$$f_{CookTorrance} = \frac{DFG}{4(\omega_o \cdot n)(\omega_i \cdot n)}$$

$$G_{Smith}(l, v) = G_{GGX}(l) \cdot G_{GGX}(v)$$

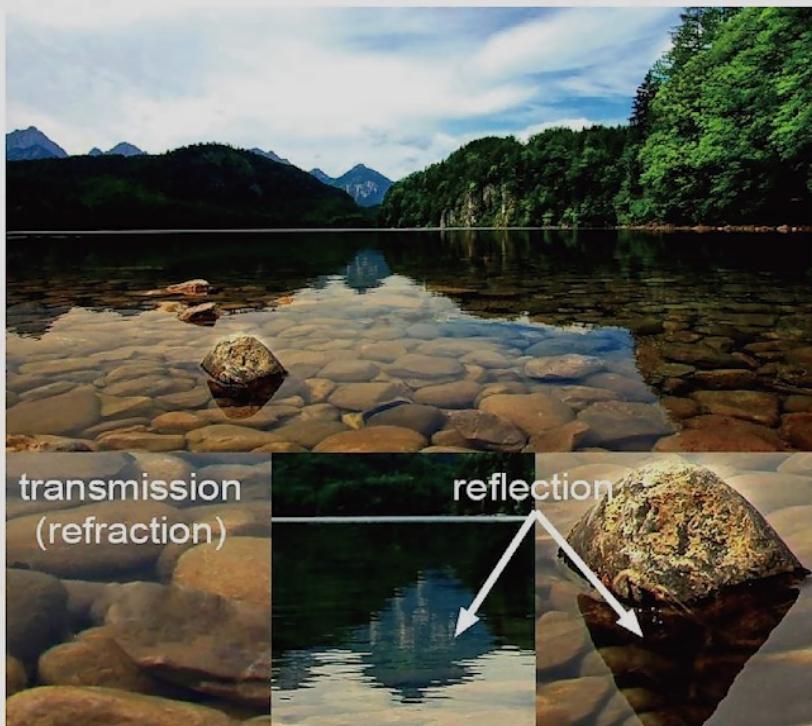
$$G_{GGX}(v) = \frac{n \cdot v}{(n \cdot v)(1 - k) + k} \quad k = \frac{(\alpha + 1)^2}{8}$$



```
// Geometry Term: Geometry masking/shadowing due to microfacets
float GGX(float NdotV, float k) {
    return NdotV / (NdotV * (1.0 - k) + k);
}
```

```
float G_Smith(float NdotV, float NdotL, float roughness)
{
    float k = pow(roughness + 1.0, 2.0) / 8.0;
    return GGX(NdotL, k) * GGX(NdotV, k);
}
```

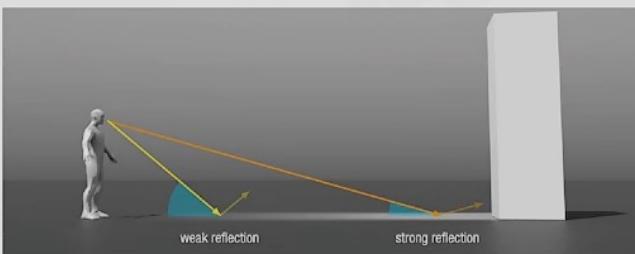
Fresnel Equation



$$f_{CookTorrance} = \frac{DFG}{4(\omega_o \cdot n)(\omega_i \cdot n)}$$

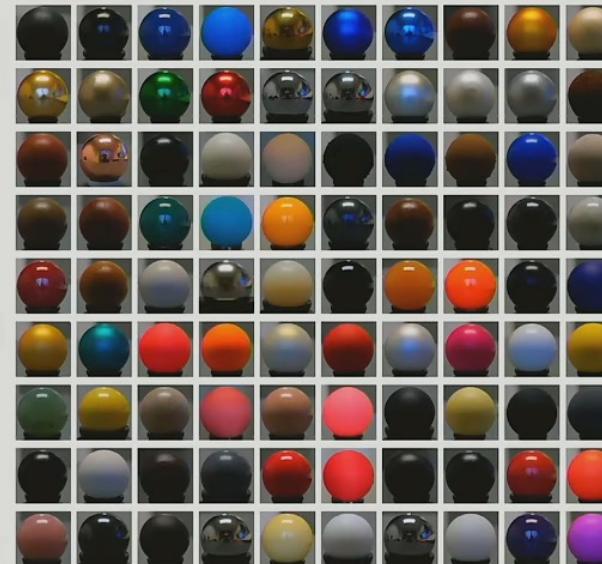
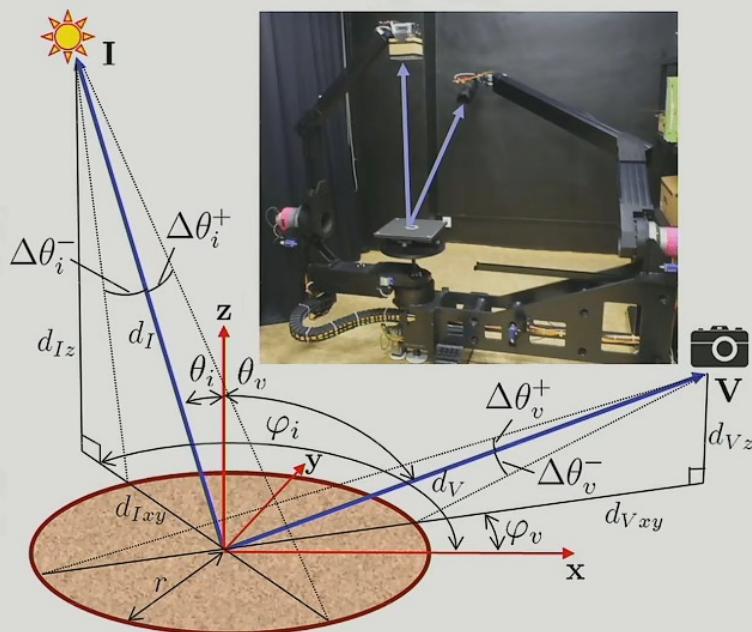
$$F_{Schlick}(h, v, F_0) = F_0 + (1 - F_0)(1 - (v \cdot h))^5$$

```
// Fresnel term with scalar optimization
float F_Schlick(float VoH, float f0)
{
    float f = pow(1.0 - VoH, 5.0);
    return f0 + (1.0 - f0) * f;
}
```



2 params.

Physical Measured Material



MERL BRDF Database of measured materials



Disney Principle Material Params.

Brent Burley

Disney Principled BRDF

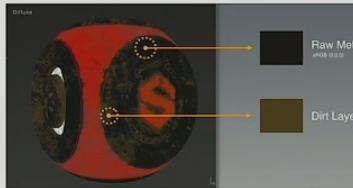
Principles to follow when implementing model:

- Intuitive rather than physical parameters should be used
- There should be as few parameters as possible
- Parameters should be zero to one over their plausible range
- Parameters should be allowed to be pushed beyond their plausible range where it makes sense
- All combinations of parameters should be as robust and plausible as possible

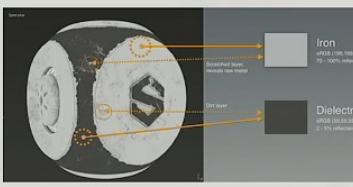
PBR Specular Glossiness



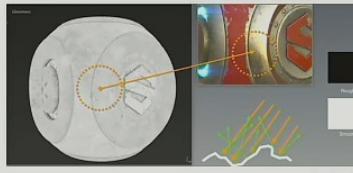
Material - SG



Diffuse - RGB - sRGB



Specular - RGB - sRGB



Glossiness - Grayscale - Linear

```
struct SpecularGlossiness
{
    float3 specular;
    float3 diffuse;
    float3 normal;
    float glossiness;
};

SpecularGlossiness getPBRParameterSG()
{
    SpecularGlossiness specular_glossiness;
    specular_glossiness.diffuse = sampleTexture(diffuse_texture, uv).rgb;
    specular_glossiness.specular = sampleTexture(specular_texture, uv).rgb;
    specular_glossiness.normal = sampleTexture(normal_texture, uv).rgb;
    specular_glossiness.glossiness = sampleTexture(gloss_texture, uv).r;
    return specular_glossiness;
}
```

```
float3 calculateBRDF(SpecularGlossiness specular_glossiness)
{
    float3 half_vector = normalize(view_direction + light_direction);
    float N_dot_L = saturate(dot(specular_glossiness.normal, light_direction));
    float N_dot_V = abs(dot(specular_glossiness.normal, view_direction));
    float3 N_dot_H = saturate(dot(specular_glossiness.normal, half_vector));
    float3 V_dot_H = saturate(dot(view_direction, half_vector));

    // diffuse
    float3 diffuse = k_d * specular_glossiness.diffuse / PI;

    // specular
    float roughness = 1.0 - specular_glossiness.glossiness;
    float3 F0 = specular_glossiness.specular;

    float D = D_GGX(N_dot_H, roughness);
    float3 F = F_Schlick(V_dot_H, F0);
    float G = G_Smith(N_dot_V, N_dot_L, roughness);
    float denominator = 4.0 * N_dot_V * N_dot_L + 0.001;

    float3 specular = (D * F * G) / denominator;

    // brdf
    return diffuse + specular;
}

void PixelShaderSG()
{
    SpecularGlossiness specular_glossiness = getPBRParameterSG();
    float3 brdf_reflection = calculateBRDF(specular_glossiness);
    return brdf_reflection * light_intensity * cos(light_incident_angle)
}
```

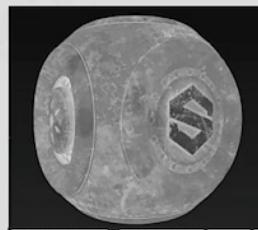
PBR Metallic Roughness



Material - MR



Base Color - RGB - sRGB



Roughness - Grayscale - Linear



Metallic - Grayscale - Linear

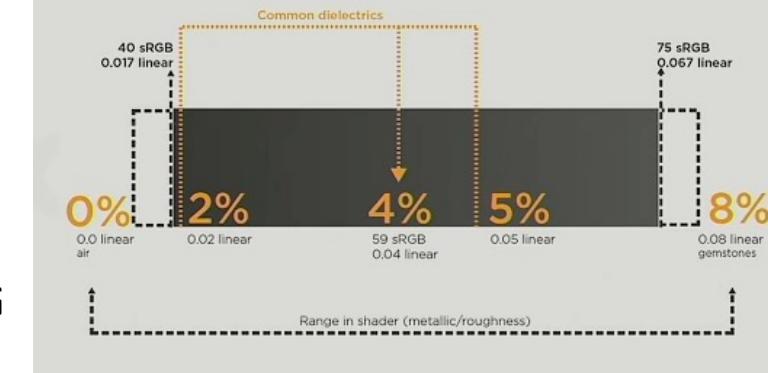
struct MetallicRoughness

```
{  
    float3 base_color;  
    float3 normal;  
    float roughness;  
    float metallic;  
};
```

```
SpecularGlossiness ConvertMetallicRoughnessToSpecularGlossiness(MetallicRoughness metallic_roughness)  
{  
    float3 base_color = metallic_roughness.base_color;  
    float roughness = metallic_roughness.roughness;  
    float metallic = metallic_roughness.metallic;  
  
    float3 dielectricSpecularColor = float3(0.08f * dielectricSpecular);  
    float3 specular = lerp(dielectricSpecularColor, base_color, metallic);  
    float3 diffuse = base_color - base_color * metallic;  
  
    SpecularGlossiness specular_glossiness;  
    specular_glossiness.specular = specular;  
    specular_glossiness.diffuse = diffuse;  
    specular_glossiness.glossiness = 1.0f - roughness;  
  
    return result;  
}
```

Dielectric	F0 (Linear)	F0 (sRGB)	Color
Water	0.02	39	
Living tissue	0.02–0.04	39–56	
Skin	0.028	47	
Eyes	0.025	44	
Hair	0.046	61	
Teeth	0.058	68	
Fabric	0.04–0.056	56–67	
Stone	0.035–0.056	53–67	
Plastics, glass	0.04–0.05	56–63	
Crystal glass	0.05–0.07	63–75	
Gems	0.05–0.08	63–80	
Diamond-like	0.13–0.2	101–124	

Convert MR to SG



PBR Pipeline MR vs SG



MR

Pros

- Can be easier to author and less prone to errors caused by supplying incorrect dielectric F0 data
- Uses less texture memory, as metallic and roughness are both grayscale maps

Cons

- No control over F0 for dielectrics in map creation. However, most implementations have a specular control to override the base 4% value
- Edge artifacts are more noticeable, especially at lower resolutions

SG

Pros

- Edge artifacts are less apparent
- Control over dielectric F0 in the specular map

Cons

- Because the specular map provides control over dielectric F0, it is more susceptible to use of incorrect values. It is possible to break the law of conservation if handled incorrectly in the shader
- Uses more texture memory with an additional RGB map

Image Based Lighting (IBL)

Basic Idea of IBL

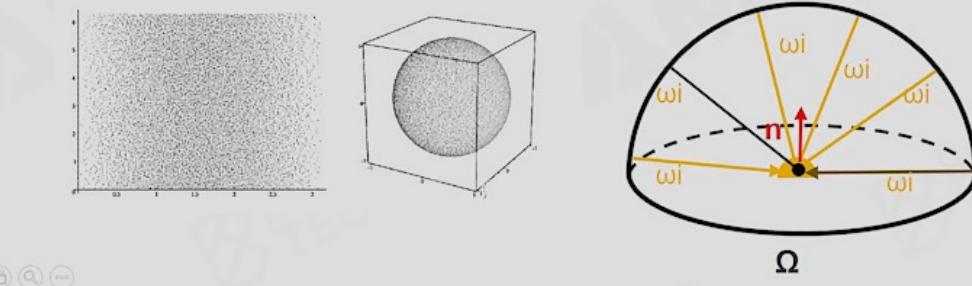
- An image representing distant lighting from all directions.



- How to shade a point under the lighting?
Solving the rendering equation:

$$L_o(\mathbf{x}, \omega_o) = \int_{H^2} f_r(\mathbf{x}, \omega_o, \omega_i) L_i(\mathbf{x}, \omega_i) \cos \theta_i d\omega_i$$

- Using Monte Carlo integration
Large amount of sampling - Slow!



Recall BRDF Function

$$L_o(\mathbf{x}, \omega_o) = \int_{H^2} f_r(\mathbf{x}, \omega_o, \omega_i) L_i(\mathbf{x}, \omega_i) \cos \theta_i d\omega_i$$

$$f_r = k_d [f_{Lambert}] + [f_{CookTorrance}]$$

diffuse specular

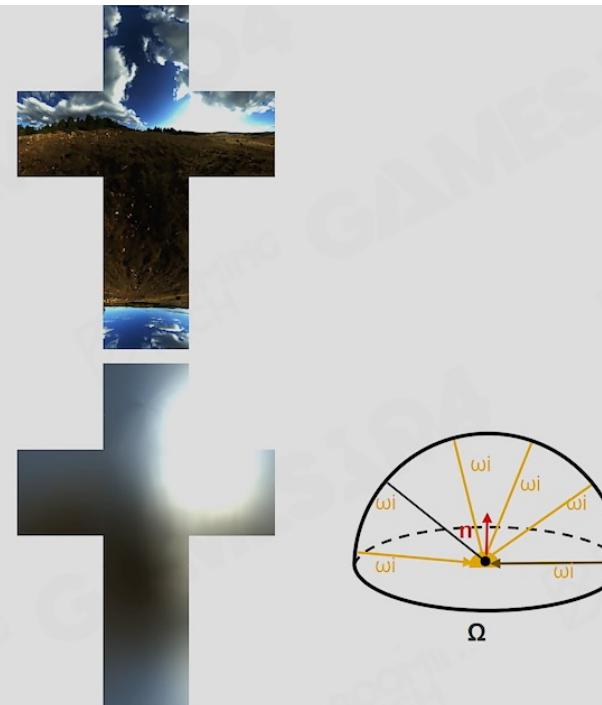
$$\begin{aligned} L_o(\mathbf{x}, \omega_o) &= \int_{H^2} (k_d f_{Lambert} + f_{CookTorrance}) L_i(\mathbf{x}, \omega_i) \cos \theta_i d\omega_i \\ &= \int_{H^2} k_d f_{Lambert} L_i(\mathbf{x}, \omega_i) \cos \theta_i d\omega_i + \int_{H^2} f_{CookTorrance} L_i(\mathbf{x}, \omega_i) \cos \theta_i d\omega_i \\ &= [L_d(\mathbf{x}, \omega_o)] + [L_s(\mathbf{x}, \omega_o)] \end{aligned}$$

Diffuse Irradiance Map

- Irradiance Map

$$\begin{aligned} f_{Lambert} &= \frac{c}{\pi} \\ \text{diffuse} & \\ L_d(\mathbf{x}, \omega_o) &= \int_{H^2} k_d f_{Lambert} L_i(\mathbf{x}, \omega_i) \cos \theta_i d\omega_i \end{aligned}$$

$$\approx k_d^* c \left[\frac{1}{\pi} \int_{H^2} L_i(\mathbf{x}, \omega_i) \cos \theta_i d\omega_i \right]$$

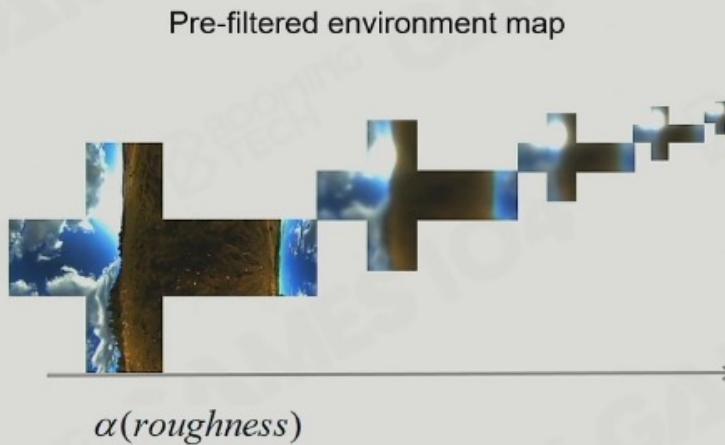


Diffuse Irradiance Map

Approximation: part (1/2)

The Lighting Term :

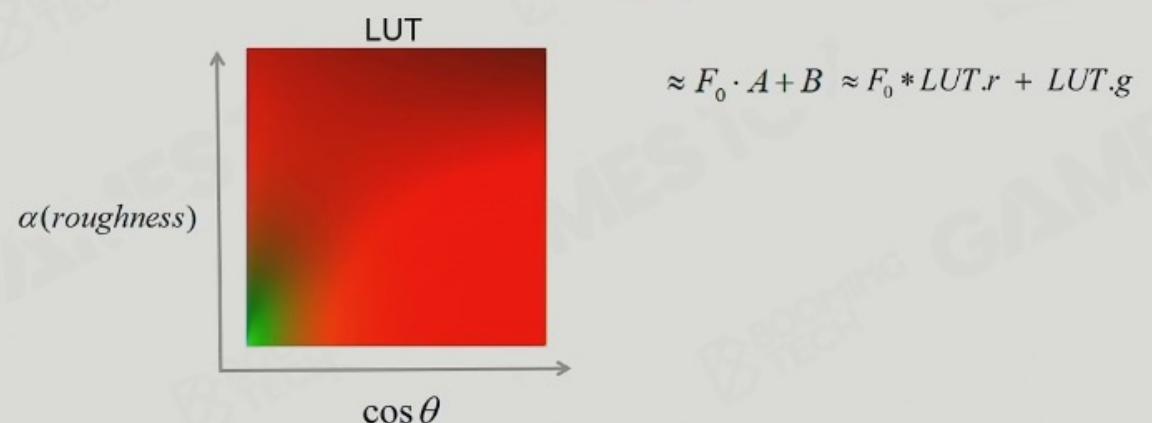
$$L_s(x, \omega_o) = \frac{\int_{H^2} f_{CookTorrances} L_i(x, \omega_i) \cos \theta_i d\omega_i}{\int_{H^2} f_{CookTorrances} \cos \theta_i d\omega_i} \cdot \int_{H^2} f_{CookTorrances} \cos \theta_i d\omega_i$$
$$\approx \frac{\sum_k^N L(\omega_i^k) G(\omega_i^k)}{\sum_k^N G(\omega_i^k)} \rightarrow \alpha$$



Approximation: part (2/2)

The BRDF Term:

$$L_s(x, \omega_o) = \frac{\int_{H^2} f_{CookTorrances} L_i(x, \omega_i) \cos \theta_i d\omega_i}{\int_{H^2} f_{CookTorrances} \cos \theta_i d\omega_i} \cdot \int_{H^2} f_{CookTorrances} \cos \theta_i d\omega_i$$
$$\approx F_0 \int_{H^2} \frac{f_{CookTorrances}}{F} (1 - (1 - \cos \theta_i)^5) \cos \theta_i d\omega_i + \int_{H^2} \frac{f_{CookTorrances}}{F} (1 - \cos \theta_i)^5 \cos \theta_i d\omega_i$$



Quick shading with precomputation

$$L(\omega_o) = k_d^* c \cdot IrranianceMap(n) + PreFiltered(R, \alpha) \cdot (F_0 * LUT.r + LUT.g)$$



Shading PBR with IBL



IBL OFF

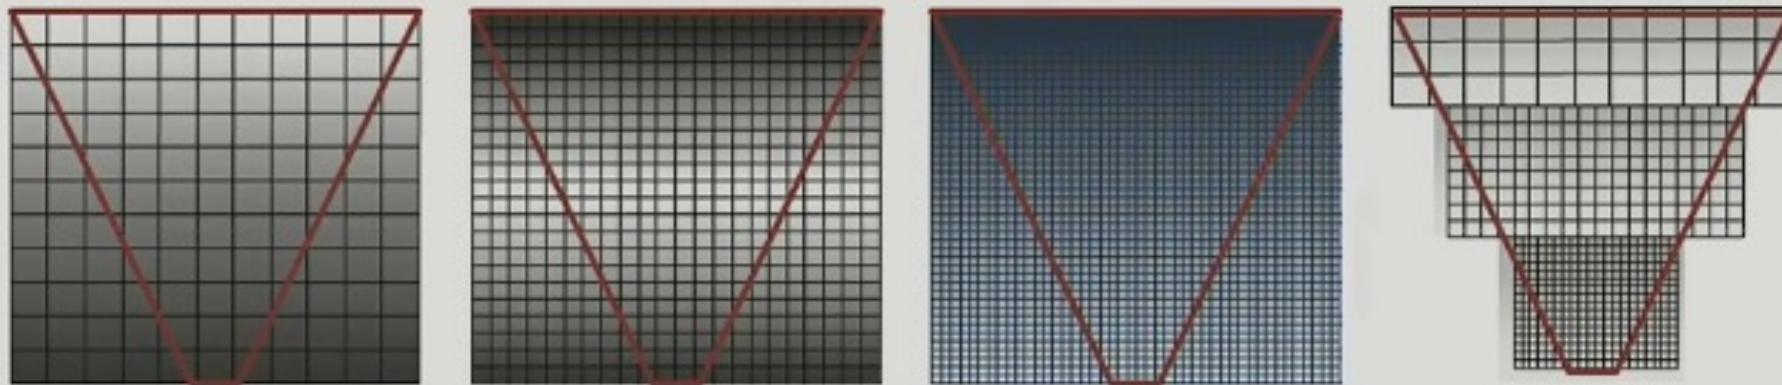


IBL ON

Shadow

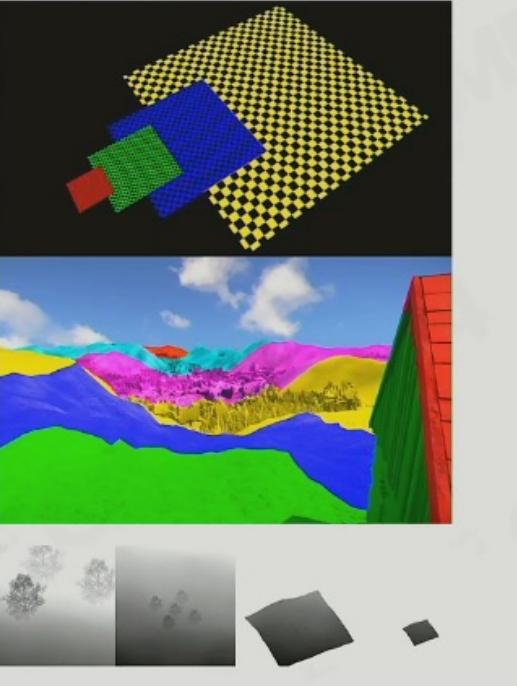
Big World and Cascade Shadow

- Partition the frustum into multiple frustums
- A shadow map is rendered for each subfrustum
- The pixel shader then samples from the map that most closely matches the required resolution



Steps of Cascade Shadow

```
splitFrustumToSubfrustums();  
calculateOrthoProjectionsForEachSubfrustum();  
renderShadowMapForEachSubfrustum();  
renderScene();  
  
vs_main()  
{  
    calculateWorldPosition()  
    ...  
}  
  
ps_main()  
{  
    transformWorldPositionsForEachProjections()  
    sampleAllShadowMaps()  
    compareDepthAndLightingPixel()  
    ...  
}
```



Cascade shadow
Different levels.

Blend between Cascade Layers

1. A visible seam can be seen where cascades overlap
2. between cascade layers because the resolution does not match.
3. The shader then linearly interpolates between the two values based on the pixel's location in the blend b



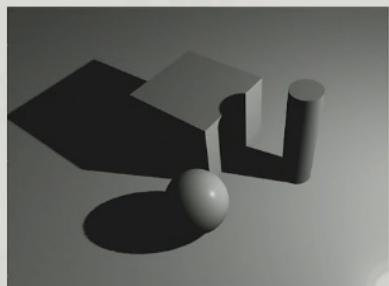
- **Pros**

- best way to prevent errors with shadowing: perspective aliasing
- fast to generate depth map, 3x up when depth writing only
- provide fairly good results

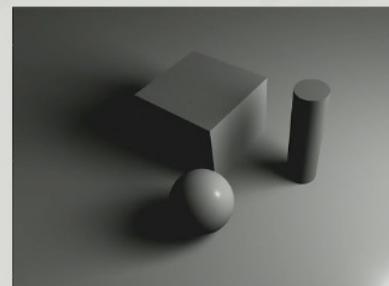
- **Cons**

- Nearly impossible to generate high quality area shadows
- No colored shadows. Translucent surfaces cast opaque shadows

Hard Shadow vs Realistic Shadow



Hard Shadow



Realistic Shadow

PCF - Percentage Closer Filter

• Target problem

- The shadows that result from shadow mapping aliasing is serious

• Basic idea

- Sample from the shadow map around the current pixel and compare its depth to all the samples
- By averaging out the results we get a smoother line between light and shadow



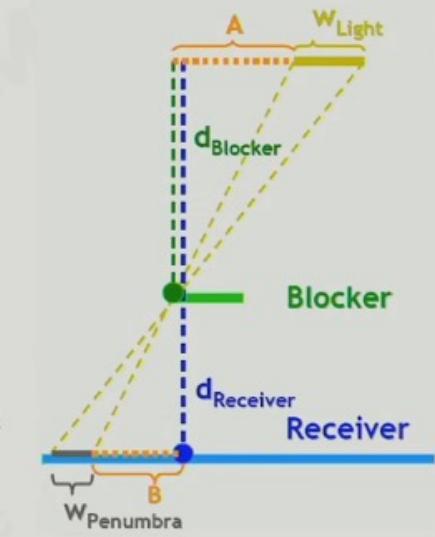
PCSS - Percentage Closer Soft Shadow

• Target problem

- Suffers from aliasing and undersampling artifacts

• Basic idea

- Search the shadow map and average the depths that are closer to the light source
- Using a parallel planes approximation



$$w_{Penumbra} = \frac{(d_{Receiver} - d_{Blocker}) \cdot w_{Light}}{d_{Blocker}}$$

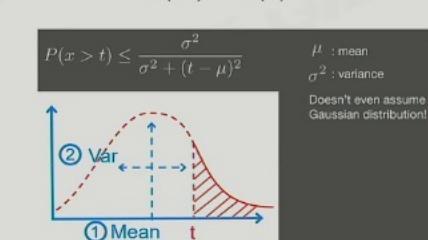
Variance Soft Shadow Map

• Target problem

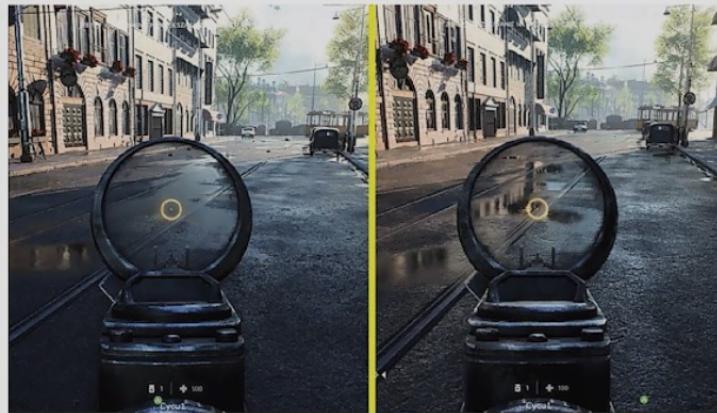
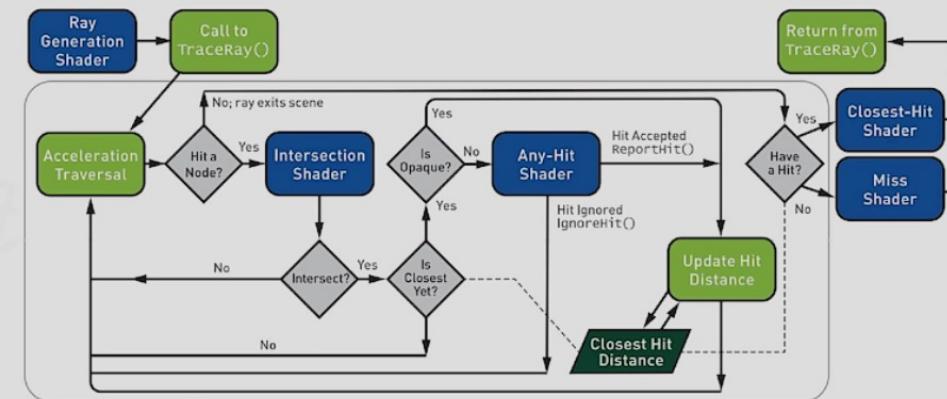
- rendering plausible soft shadow in real-time

• Basic idea

- Based on Chebyshev's inequality, using the average and variance of depth, we can approximate the percentage of depth distribution directly instead of comparing a single depth to a particular region(PCSS).

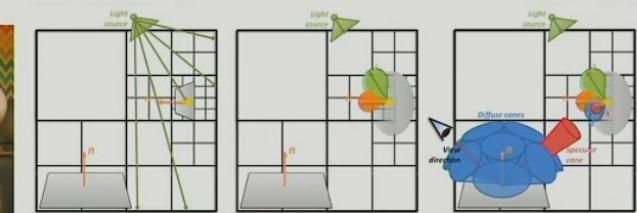


Real-Time Ray-Tracing on GPU



Real-Time Global Illumination

Screen-space GI
SDF Based GI
Voxel-Based GI (SVOGI/VXGI)
RSM / RTX GI



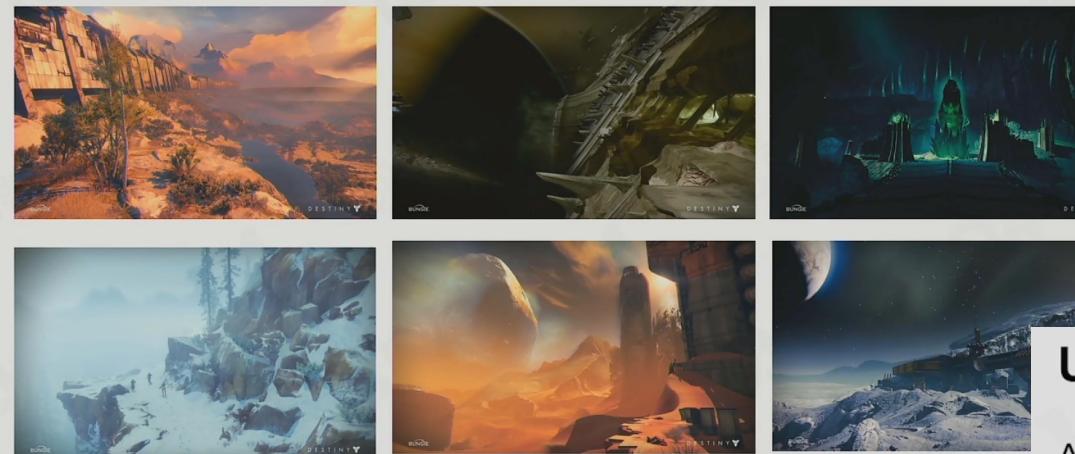
GI Off



GI On

Shader Management

Ocean of Shaders



Blow of Shaders



Uber Shader and Variants

A combination of shader for all possible light types, render passes and material types

- Shared many state and codes
- Compile to many variant short shaders by pre-defined macro

```
// sky light
#ifndef ENABLE_SKY_LIGHT
    if MATERIAL_TWOSIDED && LQ_TEXTURE_LIGHTMAP
        if (NoL == 0)
        {
        #endif

        #if MATERIAL_SHADINGMODEL_SINGLELAYERWATER
            ShadingModelContext.WaterDiffuseIndirectLuminance += SkyDiffuseLighting;
        #endif

        Color += SkyDiffuseLighting * half3(ResolvedView.SkyLightColor.rgb) *
            ShadingModelContext.DiffuseColor * MaterialAO;

    #if MATERIAL_TWOSIDED && LQ_TEXTURE_LIGHTMAP
    }
    #endif
#endif
```

Artist Create Infinite More Shaders



opaque_forward_scene_d3d11_ps_1117004100.shader.ast	2022/4/11 10:39	AST 文件	12 KB	particle_complex_particle_d3d11_vs_1117004100.shader.ast	2022/4/11 10:58	AST 文件	7 KB

165 Uber shader generated 75,536 shaders for runtime

Designing

- Entry point
 - Application layout
 - Window layout
 - Input
 - ->Events
 - Renderer
 - Render API abstraction
 - Debugging support
 - Scripting language
 - Memory systems
 - Entity-component systems (ECS)
 - Physics
 - File I/O, VFS
 - Build System
- Game engine is actually dll library
In a solution, add two projects engine(dll) and sandbox(exe)
Sandbox add reference to engine. (link)